

---

# NASH BARGAINING PREFERENCE OPTIMIZATION

**Anonymous authors**

Paper under double-blind review

## ABSTRACT

Aligning a language model with multiple objectives requires an explicit compromise between improvements over a reference policy. We introduce Nash Bargaining Preference Optimization (NBPO), which maximizes the product of objective-wise improvements in entropy-regularized preference-game values. Each objective uses a policy-dependent opponent. A constrained proximal formulation handles the zero-surplus reference boundary; its dual yields shared inverse-surplus weights. We solve the resulting finite-pool problems directly and fit the language model to their pairwise log-ratio targets. Under joint strict feasibility, the population optimum is Pareto efficient in game-value space, and exact proximal updates achieve an  $O(1/T)$  last-iterate objective gap. Controlled experiments support adaptive-game representations as nontransitivity increases. On four-objective UltraFeedback, however, the current one-stage neural results do not establish an advantage for Nash aggregation or adaptive opponents. Uniform scalarized DPO has higher judged objective means across three seeds, alongside instruction-following and safety regressions in the reported capability panel. An independent development diagnostic finds target improvements in instruction following and helpfulness, but little separation between NBPO and its exact-target controls. Neural fitting does not reproduce target distributions, yet can improve independently judged pool quality, leaving the natural-language advantage of the proposed mechanism unresolved.

## 1 INTRODUCTION

Reinforcement Learning from Human Feedback (RLHF) commonly learns scalar rewards from pairwise comparisons (Christiano et al., 2017; Ouyang et al., 2022). Direct Preference Optimization removes the separate reward model but retains a Bradley–Terry (BT) preference model (Rafailov et al., 2023; Bradley & Terry, 1952). For each prompt and objective, BT assigns one scalar score to each response and models preference probabilities through score differences. It permits noisy, contradictory individual labels, but its population preferences cannot have a cyclic strict majority relation.

Consider the prompt “How can I learn Python loops?” and three correct responses: A gives a short worked example, B a detailed conceptual explanation, and C a guided practice exercise. All annotators use the same overall-helpfulness rubric but differ in how they prefer to learn. Suppose three equally sized groups have fixed rankings: those seeking an immediate example prefer A, then B, then C; those seeking conceptual understanding prefer B, then C, then A; and those preferring practice prefer C, then A, then B. Two thirds therefore prefer A to B, two thirds B to C, and two thirds C to A. Every group’s ranking is transitive, yet the population majority cycles. This hypothetical construction illustrates how annotator heterogeneity within one rubric can produce probabilities that no shared BT reward can represent. PROSPER reports that LLM-judge comparisons on WildChecklists can remain cyclic even when each checklist item is evaluated separately (Zhang et al., 2026).

Nash Learning from Human Feedback (NLHF) retains general pairwise payoffs, and INPO, SPPO, and related methods provide scalable equilibrium-seeking updates (Munos et al., 2024; Zhang et al., 2025; Wu et al., 2025; Calandriello et al., 2024; Rosset et al., 2024; Zhou et al., 2025). Multiple objectives additionally require a compromise rule. Existing methods use scalarization, constraints, policy interpolation, preference conditioning, or Blackwell criteria (Zhou et al., 2024; Ramé et al., 2023; Zhong et al., 2024; Agnihotri et al., 2026; Zhang et al., 2026).

We use classical Nash bargaining to maximize the product of improvements over a common fallback (Nash, 1950). The bargaining parties are objective rubrics, with annotators supplying comparisons for each rubric. We give these objectives equal bargaining weight to study a compromise without externally specified priorities. Equal weight is a modeling choice: duplicating or splitting a rubric changes its influence. Prior work studies disagreement-adjusted Nash welfare over learned party rewards and bargaining over task-gradient improvements (Zhong et al., 2026; Navon et al., 2022).

Nash Bargaining Preference Optimization (NBPO) bargains over reference-relative *policy game values*. Each objective evaluates the learner against an adaptive opponent that emphasizes responses exposing its weaknesses, while a KL penalty keeps the opponent near the reference policy. The reference thus serves both as the fallback for measuring improvement and as the anchor restricting the opponent. These roles can be varied separately (Appendix A.3). Improvement concerns this adaptive game evaluation, not necessarily direct win rate against the reference; population guarantees do not imply that every prompt improves.

INPO eliminates the policy normalizer through pairwise log-ratios (Zhang et al., 2025), and PROSPER combines KL-regularized adversarial comparisons, exponential reweighting, and regression-based policy fitting (Zhang et al., 2026). NBPO adds a logarithmic bargaining objective that is undefined at the unchanged reference and couples each policy update to its adaptive opponents. We derive a constrained proximal formulation and its dual, yielding objective weights shared across prompts. At fixed weights, we directly optimize sampled-response probabilities while accounting for the opponents’ dependence on them, then fit the LLM after freezing the target. This replaces an undamped policy–opponent alternation that can be unstable even when its fixed point is optimal. Numerical solution accuracy, neural fitting error, and independent policy quality require separate checks.

Our contributions are:

- We formulate multi-objective alignment as Nash bargaining over objective-wise regularized preference-game values, separating within-objective preference representation from across-objective compromise selection.
- We derive the dual of the coupled bargaining update and implement its inner solve as small strictly concave optimization problems over sampled responses. The resulting target is fitted through pairwise log-ratio regression, with numerical accuracy distinguished from neural realization error.
- For exact population updates with a constant step size, we prove that the gap in the bargaining objective at the final iterate decreases inversely with the number of updates. Controlled experiments examine adaptive game values as nontransitivity increases, while UltraFeedback and human-labeled SAFERLHF experiments evaluate neural realization and objective-wise policy outcomes.

## 2 BACKGROUND

**General preference models.** We use contextual-bandit notation. A prompt (context)  $x$  is drawn from a distribution  $\mathcal{D}$ , a response  $y$  is an action from a finite set  $\mathcal{Y}_x$ , and  $\pi(y | x)$  is the policy. The reference policy is denoted by  $\pi_{\text{ref}}$ . For the finite-dimensional population analysis, the support of  $\mathcal{D}$  and every  $\mathcal{Y}_x$  are finite, and

$$\Pi = \prod_{x \in \text{supp}(\mathcal{D})} \Delta(\mathcal{Y}_x)$$

is the contextual policy simplex. Objective  $k \in [K]$  is represented by a pairwise preference model  $P_k(y \succ z | x) \in [0, 1]$  satisfying

$$P_k(y \succ z | x) + P_k(z \succ y | x) = 1, \quad P_k(y \succ y | x) = \frac{1}{2}. \quad (1)$$

We center each comparison at a tie,

$$\Delta_k(y, z | x) = P_k(y \succ z | x) - \frac{1}{2}, \quad \Delta_k(y, z | x) = -\Delta_k(z, y | x), \quad (2)$$

and define the expected objective- $k$  payoff of learner policy  $\pi$  against comparator policy  $\nu$  as

$$J_k(\pi, \nu) = \mathbb{E}_{x \sim \mathcal{D}, y \sim \pi(\cdot | x), z \sim \nu(\cdot | x)} \Delta_k(y, z | x). \quad (3)$$

Thus  $J_k(\pi, \nu) > 0$  means that  $\pi$  is preferred to  $\nu$  on average under objective  $k$ . This representation does not assume a scalar reward or transitive preferences.

**Bargaining problems.** A  $K$ -dimensional bargaining problem consists of a feasible value set  $\mathcal{G} \subseteq \mathbb{R}^K$  and a disagreement point  $d \in \mathcal{G}$ , the outcome obtained if no agreement is reached. For a feasible value vector  $v$ , define its objective-wise surplus over disagreement as

$$s_k(v) := v_k - d_k, \quad s(v) = (s_1(v), \dots, s_K(v)). \quad (4)$$

A value vector is individually rational when  $s(v) \geq 0$  componentwise, and strictly individually rational when  $s(v) > 0$ . On a compact, convex feasible set that is downward-closed within  $\{v : v \geq d\}$  and contains a strictly individually rational point, the classical Nash solution is

$$v^N \in \arg \max_{v \in \mathcal{G} : s(v) > 0} \sum_{k=1}^K \log s_k(v). \quad (5)$$

Equivalently, it maximizes the product of the objective-wise surpluses (Nash, 1950).

### 3 OBJECTIVE-WISE PREFERENCE GAMES

We first convert each objective’s pairwise preference model into a scalar policy value without imposing a scalar reward on individual responses. For objective  $k$ , an adaptive comparator  $\nu$  searches for responses against which the learner performs poorly. The reference policy  $\pi_{\text{ref}}$  anchors this comparator and later serves as the fallback in the bargaining problem.

**Definition 1** (Entropy-regularized game value). *For objective  $k$  and temperature  $\beta_k > 0$ , define*

$$V_{k, \beta_k}(\pi) = \min_{\nu} \{J_k(\pi, \nu) + \beta_k \mathbb{E}_{x \sim \mathcal{D}} \text{KL}(\nu(\cdot | x) \| \pi_{\text{ref}}(\cdot | x))\}. \quad (6)$$

The comparator minimizes the learner’s pairwise payoff, so it emphasizes responses that expose weaknesses of the current policy. The KL term prevents it from moving arbitrarily far from  $\pi_{\text{ref}}$ . Accordingly, a large  $\beta_k$  produces a local evaluation near the reference, whereas  $\beta_k \downarrow 0$  approaches the unregularized worst-comparator value.

For a candidate comparator response  $z$ , let

$$m_k^\pi(z | x) := \mathbb{E}_{y \sim \pi(\cdot | x)} \Delta_k(y, z | x) \quad (7)$$

be the learner’s expected margin against  $z$ . Smaller values identify stronger challenges. The inner problem in (6) has the closed-form solution

$$\nu_{k, \pi}^*(z | x) = \frac{\pi_{\text{ref}}(z | x) \exp\{-m_k^\pi(z | x)/\beta_k\}}{\mathbb{E}_{z' \sim \pi_{\text{ref}}(\cdot | x)} \exp\{-m_k^\pi(z' | x)/\beta_k\}}, \quad (8)$$

$$V_{k, \beta_k}(\pi) = -\beta_k \mathbb{E}_{x \sim \mathcal{D}} \log \mathbb{E}_{z \sim \pi_{\text{ref}}(\cdot | x)} \exp\{-m_k^\pi(z | x)/\beta_k\}. \quad (9)$$

Thus the opponent is obtained by reweighting reference responses toward those with smaller margins, and the game value is a soft minimum of these margins.

**Proposition 1** (Game-value structure). *For every  $k$  and  $\beta_k \geq 0$ ,  $V_{k, \beta_k}$  is continuous and concave in  $\pi$ . For  $\beta_k > 0$ , its policy-gradient coordinate is*

$$[\nabla_{\pi} V_{k, \beta_k}(\pi)]_{x, y} = \mathcal{D}(x) Q_k^\pi(y | x), \quad Q_k^\pi(y | x) := \mathbb{E}_{z \sim \nu_{k, \pi}^*(\cdot | x)} \Delta_k(y, z | x). \quad (10)$$

The proof is in Appendix A.1. The response-level score  $Q_k^\pi(y | x)$  is the expected margin of response  $y$  against the current objective- $k$  opponent; it therefore measures the local effect of increasing  $\pi(y | x)$ . An adaptive opponent can distinguish policies that have the same win rate against one fixed reference but different vulnerabilities to other responses. Appendix A.2 gives a finite cyclic example at the unregularized endpoint  $\beta = 0$ ; this illustration is not used in the later assumptions or proofs.

## 4 NASH BARGAINING OVER GAME-VALUE IMPROVEMENTS

We now combine the objective-wise game values using the reference policy as the common fallback. For each objective,

$$d_k := V_{k,\beta_k}(\pi_{\text{ref}}), \quad s_k(\pi) := V_{k,\beta_k}(\pi) - d_k. \quad (11)$$

The surplus  $s_k(\pi)$  is the improvement of  $\pi$  over the deployed reference under objective  $k$ .

**Assumption 1** (Joint improvement and support). *The reference  $\pi_{\text{ref}}$  has full support, and there exists a policy  $\bar{\pi}$  such that  $s_k(\bar{\pi}) > 0$  for every  $k$ .*

This is a substantive feasibility assumption. A finite-pool audit and failure outcomes are specified in Appendix A.5; the constrained formulation cannot make conflicting objectives jointly improvable. Under this condition, the NBPO population target is

$$\pi^N \in \arg \max_{\pi \in \Pi_+} F(\pi), \quad F(\pi) := \sum_{k=1}^K \log s_k(\pi), \quad \Pi_+ := \{\pi \in \Pi : s_k(\pi) > 0, \forall k\}. \quad (12)$$

Equivalently, introducing auxiliary lower bounds  $u_k > 0$  yields

$$(\pi^N, u^N) \in \arg \max_{\pi, u \in \mathbb{R}_{++}^K} \sum_{k=1}^K \log u_k \quad \text{s.t.} \quad u_k \leq s_k(\pi), \quad k \in [K], \quad \pi \in \Pi. \quad (13)$$

Because the objective is increasing in every  $u_k$ , each constraint is tight at an optimum and  $u_k^N = s_k(\pi^N)$ .

To connect (12) to the classical bargaining problem, define

$$V_{\beta}(\pi) = (V_{1,\beta_1}(\pi), \dots, V_{K,\beta_K}(\pi)), \quad d = (d_1, \dots, d_K),$$

and the downward-closed feasible value set

$$\mathcal{G}_{\beta} := \{v \in \mathbb{R}^K : d \leq v \leq V_{\beta}(\pi) \text{ componentwise for some } \pi \in \Pi\}. \quad (14)$$

**Proposition 2** (Policy-induced Nash bargaining). *Under Assumption 1,  $\mathcal{G}_{\beta}$  is compact, convex, downward closed above  $d$ , and nondegenerate. The policies solving (12) realize the classical Nash solution of  $(\mathcal{G}_{\beta}, d)$ . Consequently, every optimizer has positive surplus in every objective and is Pareto optimal in game-value space; all optimizers induce the same game-value vector; and the optimizer set is unchanged by independent transformations*

$$(V_{k,\beta_k}, d_k) \mapsto (c_k V_{k,\beta_k} + b_k, c_k d_k + b_k), \quad c_k > 0.$$

Appendix A.3 verifies the feasible-set conditions. Once those conditions hold, the stated outcome properties are standard consequences of the classical Nash bargaining solution (Nash, 1950); we do not claim them as new bargaining axioms. These properties concern population game values. Scaling primitive margins also requires scaling  $\beta_k$ ; arbitrary recalibration of judge probabilities is not covered (Appendix A.3). Appendix A.4 compares these guarantees with utilitarian and maxmin aggregation.

## 5 NASH BARGAINING PREFERENCE OPTIMIZATION

Optimizing (12) presents three practical difficulties. First, the reference has  $s_k(\pi_{\text{ref}}) = 0$ , so the logarithmic objective is undefined at initialization. Second, every game value contains an opponent that depends on the learner. Third, the population policy update must be realized by an autoregressive model. We handle these issues with a constrained proximal problem, a prompt-separable finite-pool inner solve, and neural log-ratio projection.

### 5.1 CONSTRAINED PROXIMAL UPDATE AND ADAPTIVE OBJECTIVE WEIGHTS

For  $\pi \in \Pi_+$ , Proposition 1 and the chain rule give

$$[\nabla_{\pi} F(\pi)]_{x,y} = \mathcal{D}(x) \sum_{k=1}^K \frac{Q_k^{\pi}(y | x)}{s_k(\pi)}. \quad (15)$$

The inverse-surplus factors emphasize objectives that have improved less. Because this expression cannot be evaluated at  $\pi_{\text{ref}}$ , outer stage  $t$  instead solves

$$\begin{aligned} (\pi_{t+1}, u_{t+1}) \in \arg \max_{\pi, u \in \mathbb{R}_{++}^K} \quad & \sum_{k=1}^K \log u_k - \frac{1}{\eta_t} D(\pi \| \pi_t) \\ \text{s.t.} \quad & u_k \leq s_k(\pi), \quad k \in [K], \\ & \pi \in \Pi, \end{aligned} \quad (16)$$

where  $\pi_0 = \pi_{\text{ref}}$ ,  $\eta_t > 0$ , and  $D(\pi \| \pi_t) = \mathbb{E}_{x \sim \mathcal{D}} \text{KL}(\pi(\cdot | x) \| \pi_t(\cdot | x))$ .

The dual is

$$\min_{\lambda \in \mathbb{R}_{++}^K} \phi_t(\lambda), \quad \phi_t(\lambda) = - \sum_{k=1}^K \log \lambda_k - K + \max_{\pi \in \Pi} \left\{ \sum_{k=1}^K \lambda_k s_k(\pi) - \frac{1}{\eta_t} D(\pi \| \pi_t) \right\}, \quad (17)$$

where  $\lambda = (\lambda_1, \dots, \lambda_K) \in \mathbb{R}_{++}^K$  is the vector of Lagrange multipliers. For fixed  $\lambda$ , let

$$\pi_t^*(\lambda) := \arg \max_{\pi \in \Pi} \left\{ \sum_{k=1}^K \lambda_k s_k(\pi) - \frac{1}{\eta_t} D(\pi \| \pi_t) \right\}. \quad (18)$$

The KL term makes the maximizing policy  $\pi_t^*(\lambda)$  unique. Danskin’s theorem (Danskin, 1967) then allows us to differentiate  $\phi_t(\lambda)$  while holding this maximizing policy fixed, yielding

$$\nabla_k \phi_t(\lambda) = s_k(\pi_t^*(\lambda)) - \frac{1}{\lambda_k}. \quad (19)$$

The complete dual derivation is in Appendix A.5.

**Proposition 3** (Inverse-surplus objective weights). *At an optimal primal–dual solution of (16),*

$$u_{t+1,k} = s_k(\pi_{t+1}), \quad \lambda_{t+1,k} = \frac{1}{u_{t+1,k}} = \frac{1}{s_k(\pi_{t+1})}. \quad (20)$$

Thus the constrained problem recovers the inverse-surplus weighting in (15) without selecting objective weights in advance.

## 5.2 DIRECT FINITE-POOL INNER SOLVE AND NEURAL PROJECTION

At outer stage  $t$ , draw learner occurrences  $\mathcal{Y}_x = \{y_{x,i}\}_{i=1}^{n_x}$  from  $\pi_t$  and comparator occurrences  $\mathcal{Z}_x = \{z_{x,j}\}_{j=1}^{m_x}$  from  $\pi_{\text{ref}}$ . Write  $\hat{A}_{k,x}[i, j] = \hat{\Delta}_k(y_{x,i}, z_{x,j} | x)$ . Retaining duplicate draws with multiplicity gives empirical base measures  $p_{t,x}(i) = 1/n_x$  and  $\mu_x(j) = 1/m_x$ . These are importance-sampling occurrence weights, not normalized LM likelihoods. For  $p \in \Delta^{n_x}$ , define

$$\hat{m}_{k,x}(p)_j := \sum_{i=1}^{n_x} p_i \hat{A}_{k,x}[i, j], \quad (21)$$

$$\hat{V}_{k,x}(p) := -\beta_k \log \sum_{j=1}^{m_x} \mu_{x,j} \exp\{-\hat{m}_{k,x}(p)_j / \beta_k\}. \quad (22)$$

The fixed disagreement estimate  $\hat{d}_k$  comes from a reference–reference tensor, not an assumed zero. Let  $\hat{\mathcal{D}}$  be the empirical prompt distribution and set  $\hat{s}_k(p) = \mathbb{E}_{x \sim \hat{\mathcal{D}}} \hat{V}_{k,x}(p_x) - \hat{d}_k$ . The shared multipliers aggregate these prompt averages; there is no separate Nash bargain at each prompt.

For fixed  $\lambda$ , the empirical counterpart of (18) separates across prompts:

$$p_{t,x}^*(\lambda) \in \arg \max_{p \in \Delta^{n_x}} \left\{ \sum_{k=1}^K \lambda_k \hat{V}_{k,x}(p) - \frac{1}{\eta_t} \text{KL}(p \| p_{t,x}) \right\}. \quad (23)$$

Each problem is strictly concave because of the KL term. The opponent is analytically eliminated inside  $\hat{V}$ , so direct optimization accounts for its dependence on  $p$ . We solve these small programs to

---

**Algorithm 1** Nash Bargaining Preference Optimization
 

---

**Require:** Reference  $\pi_{\text{ref}}$ , pairwise models  $P_{1:K}$ , temperatures  $\beta_{1:K}$ , outer stages  $T$ .

- 1: Set  $\pi_0 \leftarrow \pi_{\text{ref}}$ ; construct fixed reference pools and  $\hat{d}_{1:K}$ .
  - 2: **for**  $t = 0, \dots, T - 1$  **do**
  - 3:   Draw learner pools from  $\pi_t$ , construct cross tensors, initialize  $\lambda > 0$ .
  - 4:   **repeat**
  - 5:     Solve (23) independently for every prompt.
  - 6:     Update shared  $\lambda$  using  $\hat{s}_k - 1/\lambda_k$ ; check residuals and bounds.
  - 7:     **until** unprojected KKT tolerance is met or the numerical budget is exhausted
  - 8:     On uncertified output, audit feasibility, retain  $\pi_t$ , and return its failure status.
  - 9:     Construct  $\hat{h}_t^*$  by (24); verify (27).
  - 10:    Fit one neural candidate  $\tilde{\pi}_{t+1}$  using (26).
  - 11:    Measure post-fit diagnostics on development data; accept only if declared checks pass.
  - 12:    Set  $\pi_{t+1} \leftarrow \tilde{\pi}_{t+1}$  if accepted, otherwise retain  $\pi_t$  and record rejection.
  - 13: **end for**
  - 14: **return**  $\pi_T$ .
- 

numerical tolerance, then solve the shared dual condition  $\hat{s}_k(p_t^*(\lambda)) = 1/\lambda_k$  in  $\log \lambda$ . A computational box  $\Lambda = [\lambda_{\min}, \lambda_{\max}]^K$  requires reporting both unprojected stationarity and bound activity; projected stationarity alone does not certify the original problem. Appendix A.5 gives the residual checks and an explicit instability example for undamped policy–opponent alternation.

Once the dual converges, the finite-pool target defines the pairwise log-ratio change

$$\hat{h}_t^*(x, i, i') := \log \frac{p_{t,x}^*(i)}{p_{t,x}(i)} - \log \frac{p_{t,x}^*(i')}{p_{t,x}(i')}. \quad (24)$$

For a neural policy  $\pi_\theta$ , let

$$h_t(\pi_\theta; x, y, y') := \log \frac{\pi_\theta(y \mid x)}{\pi_t(y \mid x)} - \log \frac{\pi_\theta(y' \mid x)}{\pi_t(y' \mid x)}. \quad (25)$$

We project the finite-pool target into the LLM with

$$\mathcal{L}_{\text{reg}}^{(t)}(\theta) = \mathbb{E}_{\substack{x \sim \hat{\mathcal{D}} \\ (i, i') \sim \xi_x}} \left[ \left( h_t(\pi_\theta; x, y_{x,i}, y_{x,i'}) - \hat{h}_t^*(x, i, i') \right)^2 \right]. \quad (26)$$

The target is frozen during fitting. Complete pair coverage identifies only the reweighted sampled support, not unseen responses. At the inner optimum,

$$\hat{h}_t^*(x, i, i') = \eta_t \sum_{k=1}^K \lambda_k \left( \hat{Q}_{k,x}^*(i) - \hat{Q}_{k,x}^*(i') \right), \quad (27)$$

where  $\hat{Q}_{k,x}^* = \hat{A}_{k,x} \nu_{k,x}(p_{t,x}^*)$ . This identity checks the solver target independently. Pairwise probabilities may come from direct human comparisons or a general pairwise preference model; NBPO never requires them to decompose into scalar response rewards.

Neural pool diagnostics use  $\tilde{p}_{\theta,x}(i) \propto p_{t,x}(i) \pi_\theta(y_{x,i} \mid x) / \pi_t(y_{x,i} \mid x)$ , not normalized raw LM likelihoods. Appendix F compares regression with weighted behavior cloning of the same target.

Acceptance is empirical, uses development rather than final-test data, and provides no population certificate. Appendix A.5 defines failure reporting. Historical runs require artifact verification of this execution contract; all reported LLM runs use  $T = 1$ .

**Last-iterate convergence.** The following Bregman proximal-point guarantee concerns exact population solves, not finite-pool fitting or empirical acceptance;  $T$  counts exact solves rather than neural updates. Under Assumption 1, with  $\pi_0 = \pi_{\text{ref}}$  and (16) solved exactly at every outer iteration with arbitrary  $\eta_t > 0$ , the Nash welfare is monotone,  $F(\pi_{t+1}) \geq F(\pi_t)$  for  $t \geq 1$ , and  $F(\pi^*) - F(\pi_T) \leq D(\pi^* \parallel \pi_{\text{ref}}) / \sum_{t < T} \eta_t$ : a constant proximal step gives an  $O(1/T)$  final-iterate gap. Theorem 1 and its proof are in Appendix A.6.

Table 1: UF-4 reference wins on 1,236 common prompts, averaging orders and counting ties as 0.5. Attributes are seed means;  $\bar{W}_{\min} = S^{-1} \sum_s \min_k W_{k,s}$  averages seed-specific minima. Sample SD appears below each row. Bold denotes a reference-win interval above 0.5 at every evaluated seed; prompt uncertainty is separate (Appendix G).

| Method                            | Seeds | IF            | Truth         | Honesty | Help          | $\bar{W}_{\min}$ |
|-----------------------------------|-------|---------------|---------------|---------|---------------|------------------|
| Base (defined reference)          | –     | 0.500         | 0.500         | 0.500   | 0.500         | 0.500            |
| <i>Matched mechanism controls</i> |       |               |               |         |               |                  |
| NBPO                              | 3     | 0.5082        | 0.5087        | 0.4955  | <b>0.5478</b> | 0.4955           |
| sample SD over seeds              |       | 0.0085        | 0.0039        | 0.0017  | 0.0039        | 0.0017           |
| Fixed-reference Nash              | 3     | 0.5118        | 0.5134        | 0.5002  | <b>0.5485</b> | 0.5002           |
| sample SD over seeds              |       | 0.0081        | 0.0016        | 0.0058  | 0.0026        | 0.0058           |
| BT-RM–Nash                        | 3     | 0.5056        | 0.5131        | 0.5032  | <b>0.5466</b> | 0.5021           |
| sample SD over seeds              |       | 0.0109        | 0.0047        | 0.0061  | 0.0060        | 0.0078           |
| Game-utilitarian                  | 3     | 0.5115        | 0.5076        | 0.4997  | <b>0.5508</b> | 0.4980           |
| sample SD over seeds              |       | 0.0043        | 0.0079        | 0.0097  | 0.0055        | 0.0091           |
| Global game-maxmin                | 3     | 0.5130        | 0.5136        | 0.4944  | <b>0.5500</b> | 0.4944           |
| sample SD over seeds              |       | 0.0140        | 0.0077        | 0.0066  | 0.0050        | 0.0066           |
| <i>External baselines</i>         |       |               |               |         |               |                  |
| Scalarized DPO (uniform)          | 3     | <b>0.5653</b> | <b>0.5647</b> | 0.5044  | <b>0.6282</b> | 0.5044           |
| sample SD over seeds              |       | 0.0065        | 0.0018        | 0.0042  | 0.0017        | 0.0042           |
| PROSPER (adapt.)                  | 2     | 0.5109        | 0.5095        | 0.5031  | <b>0.5452</b> | 0.4990           |
| sample SD over seeds              |       | 0.0303        | 0.0066        | 0.0076  | 0.0070        | 0.0134           |
| MOPO (adapt.)                     | 3     | <b>0.5205</b> | 0.5171        | 0.5028  | <b>0.5278</b> | 0.5028           |
| sample SD over seeds              |       | 0.0019        | 0.0072        | 0.0029  | 0.0117        | 0.0029           |

## 6 EXPERIMENTS

We test generated-policy quality, bargaining trade-offs, and adaptive opponents. UltraFeedback (UF-4) provides score-induced four-objective supervision; SafeRLHF provides human pairwise feedback. Controlled games isolate nontransitivity. UF labels are transitive by construction, and a separate 60,000-judgment audit finds no significant repeated-cycle excess (Appendix D.1); neither natural panel currently establishes an advantage from cyclic feedback.

### 6.1 GENERATED-POLICY QUALITY UNDER SCORE-INDUCED SUPERVISION

**Protocol and dataset role.** The panel uses Llama-3.1-8B-Instruct and UltraFeedback instruction-following (IF), truthfulness, honesty, and helpfulness labels (Cui et al., 2024). Filtering leaves 38,240 eligible prompt groups, supporting disjoint 20k/2k preference-model train/dev, 10k/1k policy train/dev, and 2k final splits. Policies train on newly sampled responses for 1250 updates. UF tests multi-objective alignment and aggregation under a common training protocol; it does not test learning from cyclic labels. Transitivity within each objective does not remove conflict across objectives. Appendix D.3 gives the data, teacher, and training details.

**Matched comparisons.** Fixed-reference Nash changes opponent adaptation; BT-RM–Nash changes preference representation; game-utilitarian and global game-maxmin change aggregation over the same prompt-averaged game values. External rows use uniform scalarized DPO and PROSPER/MOPO adaptations, whose implementation contracts are in Appendix D.3. PROSPER uses prompt-dependent aggregation; the matched maxmin control uses global aggregation. The MOPO row is a plug-in adaptation, not the full published robustified method. Matching coefficient L1 norms does not match realized policy KL.

**Evaluation and current results.** All policies use the same planned 2000 final prompts and cached base responses. A prompt enters Table 1 only if every reported arm has both orders parsed for all four attributes; results are recomputed when this intersection changes. Whole-prompt bootstrap replicates recompute the minimum objective, with prompt uncertainty kept separate from policy-seed variation. No reported minimum-objective interval excludes 0.5, and these results do not establish a Nash advantage.




Figure 1: Controlled nontransitivity: normalized worst surplus versus BT deviance. Points and bars show means and one sample standard deviation over five seeds; all six plotted methods converge.

## 6.2 BARGAINING AND EMPIRICAL TRADE-OFFS

The matched mechanisms yield similar reference-win estimates, without a resolved advantage for Nash aggregation or adaptive opponents (Table 1); this does not establish equivalence. Uniform DPO exceeds each matched mechanism’s mean on all four attributes, with NBPO gaps of 0.0571, 0.0560, 0.0090 and 0.0804. Its capability and safety losses motivate comparison at matched policy movement or capability retention (Section B.1). PROSPER and MOPO produce different trade-offs, with substantial PROSPER seed variation; neither establishes superiority across all outcomes.

Nash and utilitarian targets are close: mean TV is 0.0195 on UF-4 and 0.0008 on SafeRLHF (Appendix F). Weak target separation and neural realization error are examined in Appendix B.

Appendix Figure 2 projects the four-objective outcomes. The declared seven-setting DPO sweep has reconciled uniform and two specialist settings: the IF specialist reaches .5845 IF but .4775 honesty, and the truthfulness specialist reaches .5819 truthfulness and .5241 honesty. The specialist points each use one seed; the sweep remains incomplete. Operating points are selected on development data, and dominance requires the full four-dimensional vector (Appendix G.4).

## 6.3 ADAPTIVE OPPONENTS AND PREFERENCE REPRESENTATION

**Controlled nontransitivity.** We perturb transitive games while preserving a positive joint-improvement margin and evaluate all methods with common adaptive-game values. At the strongest perturbation, normalized worst surplus is 0.931 for NBPO,  $-0.241$  for fixed-reference Nash, and  $-0.242$  for BT-RM-Nash (Figure 1). Game-utilitarian (0.908) and Game-KS (0.932) also perform well. This supports the adaptive-game representation on these constructions without establishing a unique Nash advantage. A second construction gives the same trend with smaller effects (Appendix C). The natural UF-4 results do not establish this representation advantage.

**Independent natural-policy cross-play.** Fresh held-out responses compare the base, NBPO, fixed-reference Nash, and a competitor selected on development data. For independently judged, order-averaged wins  $W_k(\pi, q)$  and a frozen comparator bank  $\mathcal{B}$ , we report

$$W_{\text{ref}}^{\min}(\pi) = \min_k W_k(\pi, \pi_{\text{ref}}), \quad W_{\mathcal{B}}^{\min}(\pi) = \min_{k, q \in \mathcal{B}} W_k(\pi, q).$$

Worst-bank wins measure vulnerability to these comparators, not population exploitability. Table 2 preserves the supplied summary, whose valid prompt counts differ across rows and statistics; a common-set recomputation is required for paired method comparisons. Appendix G.2 retains the objective-wise matrices and regularized bank diagnostic. A self-comparison fixed at 0.5 bounds  $W_{\mathcal{B}}^{\min}$  above by 0.5 for bank members, so reaching 0.5 alone does not establish robustness beyond the bank.

**Capabilities and target realization.** NBPO’s IFEval score is 0.7264 versus the base’s 0.7542, with harmful-completion rate 0.3906 versus 0.2812; scalarized DPO also incurs capability regressions (Table 4). The independent development diagnostic finds target-versus-source gains in IF and helpfulness,

Table 2: One-seed cross-play on 500 planned prompts. Reference counts are row-specific; bank counts span 446–480, with 417 prompts common to all pairs and rubrics. These unequal-denominator summaries are descriptive. All trained-policy surplus intervals include zero.

| Policy                                 | Seeds | $N_{\text{ref}}$ | $W_{\text{ref}}^{\min}$ | $W_{\mathcal{B}}^{\min}$ | $\min_k \hat{s}_k^{\mathcal{B}}$ |
|----------------------------------------|-------|------------------|-------------------------|--------------------------|----------------------------------|
| Base                                   | –     | 446              | 0.5000                  | 0.4439                   | +0.0000                          |
| NBPO                                   | 1     | 477              | 0.5068                  | 0.5000                   | +0.0149                          |
| Fixed-reference Nash                   | 1     | 471              | 0.4904                  | 0.4771                   | −0.0133                          |
| Dev-selected competitor (util_mse_s42) | 1     | 480              | 0.4990                  | 0.4897                   | +0.0106                          |

but NBPO’s exact target differs from fixed-reference and utilitarian targets by at most 0.0018 in direct wins. Fitted pools improve several independently judged scores despite imperfect target realization; the projection pilot does not resolve a fresh-policy advantage. Detailed coverage, target/fitted/fresh comparisons and supporting SafeRLHF results are in Appendix B. These observations motivate the prospective fixed-objective conflict and cyclicity tests in Appendix I; their results are reported separately from the completed UF panel.

## 7 RELATED WORK

**General pairwise preferences.** NLHF and its scalable variants optimize directly against general pairwise preference models rather than fitting scalar rewards (Munos et al., 2024; Wu et al., 2025; Zhang et al., 2025; Calandriello et al., 2024; Rosset et al., 2024; Zhou et al., 2025). Blackwell-style methods extend this view to vector-valued feedback and target sets (Blackwell, 1956; Bhatia et al., 2020). PROSPER introduces a maximum-entropy Blackwell winner with a scalable single-policy update and an inner worst-case choice over objective weights and opponents (Zhang et al., 2026). NBPO instead keeps one regularized game value per objective and applies Nash bargaining across the improvements over a shared fallback.

**Multi-objective alignment.** MODPO, Rewarded Soups, and Panacea expose trade-offs through scalarization, policy interpolation, or preference conditioning (Zhou et al., 2024; Ramé et al., 2023; Zhong et al., 2024). MaxMin-RLHF and RMOD optimize weakest-objective criteria (Chakraborty et al., 2024; Son et al., 2026). MOPO uses pairwise preferences without point-wise rewards, selecting a primary objective and lower-bound constraints with a lagged reference in its practical algorithm (Agnihotri et al., 2026). RACO resolves conflicts between DPO-style objective gradients (Chen et al., 2026). These approaches differ from bargaining over reference-relative robust policy values.

**Social choice and bargaining.** The Nash solution and its Pareto, uniqueness, and affine-invariance properties are classical (Nash, 1950). Social-choice work argues that alignment should state its aggregation rule explicitly (Conitzer et al., 2024). Multi-party RLHF studies utilitarian, Nash, and leximin welfare over party-specific learned rewards and separately considers reward-free von Neumann solutions (Zhong et al., 2026); Nash bargaining also selects immediate task-gradient improvements in multi-task learning (Navon et al., 2022). Our contribution is not a new bargaining axiom, but a formulation and practical optimizer for applying the classical Nash rule to policy-dependent objective-wise preference-game values.

## 8 CONCLUSION AND LIMITATIONS

NBPO combines reference-surplus bargaining with adaptive preference-game evaluation and a direct finite-pool proximal solver. Controlled experiments support the adaptive representation; the current one-stage natural-language results do not establish a Nash advantage. The development diagnostic finds weak target separation and imperfect neural realization; their contributions to fresh-response differences remain unresolved. Joint improvement may be infeasible, and exact population guarantees do not certify a fitted neural policy, direct reference wins or per-prompt safety.

---

## AI USE STATEMENT

In this work, generative AI tools were used to critique the problem formulation, assist with mathematical derivations and proof presentation, propose experimental checks, and edit manuscript text. The authors are responsible for independently verifying every theorem, citation, implementation, empirical result, and disclosure in the final submission.

## ETHICS STATEMENT

This work studies aggregation rules for language-model alignment, including harmlessness as one objective. Improvement in a game value is not an absolute safety certificate. Safety-critical deployment should retain hard constraints and independent safety evaluations in addition to the bargaining objective.

## REPRODUCIBILITY STATEMENT

Algorithm 1 distinguishes the population proximal problem, the numerically solved finite-pool inner problem, and its neural projection. Appendix A contains the dual and convergence proofs; Appendices C–G specify constructions, supervision, hyperparameters, diagnostics, and evaluation protocols. The submission package should include candidate pools, preference tensors, dual traces, solver residuals, per-seed outputs, and the code and configurations required to reproduce the reported results.

## REFERENCES

- Akhil Agnihotri, Rahul Jain, Deepak Ramachandran, and Zheng Wen. Multi-objective preference optimization: Improving human alignment of generative models. In *International Conference on Machine Learning*, 2026.
- Kush Bhatia, Ashwin Pananjady, Peter L. Bartlett, Anca D. Dragan, and Martin J. Wainwright. Preference learning along multiple criteria: A game-theoretic perspective. In *Advances in Neural Information Processing Systems*, volume 33, pp. 7413–7424, 2020.
- David Blackwell. An analog of the minimax theorem for vector payoffs. *Pacific Journal of Mathematics*, 6(1):1–8, 1956.
- Ralph Allan Bradley and Milton E. Terry. Rank analysis of incomplete block designs: I. the method of paired comparisons. *Biometrika*, 39(3/4):324–345, 1952.
- Daniele Calandriello, Daniel Guo, Rémi Munos, Mark Rowland, Yunhao Tang, Bernardo Avila Pires, Pierre Harvey Richemond, Charline Le Lan, Michal Valko, Tianqi Liu, et al. Human alignment of large language models through online preference optimisation. *arXiv preprint arXiv:2403.08635*, 2024.
- Souradip Chakraborty, Jiahao Qiu, Hui Yuan, Alec Koppel, Furong Huang, Dinesh Manocha, Amrit Singh Bedi, and Mengdi Wang. Maxmin-rlhf: Alignment with diverse human preferences. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, 2024.
- Peter Chen, Xiaopeng Li, Xi Chen, and Tianyi Lin. Reward-free alignment for conflicting objectives. *arXiv preprint arXiv:2602.02495*, 2026. Accepted to the 43rd International Conference on Machine Learning (Oral).
- Paul F. Christiano, Jan Leike, Tom Brown, Miljan Martic, Shane Legg, and Dario Amodei. Deep reinforcement learning from human preferences. In *Advances in Neural Information Processing Systems*, volume 30, 2017.
- Vincent Conitzer, Rachel Freedman, Jobst Heitzig, Wesley H. Holliday, Bob M. Jacobs, Nathan Lambert, Milan Mossé, Eric Pacuit, Stuart Russell, Hailey Schoelkopf, Emanuel Tewolde, and William S. Zwicker. Position: Social choice should guide AI alignment in dealing with diverse human feedback. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, 2024.

- Ganqu Cui, Lifan Yuan, Ning Ding, Guanming Yao, Bingxiang He, Wei Zhu, Yuan Ni, Guotong Xie, Yankai Lin, and Maosong Sun. Ultrafeedback: Boosting language models with scaled ai feedback. In *International Conference on Machine Learning (ICML)*, 2024.
- John M. Danskin. *The Theory of Max-Min and Its Application to Weapons Allocation Problems*. Springer-Verlag, 1967.
- Jiwoo Han, Moulinath Banerjee, and Yuekai Sun. Maximin relative improvement: Fair learning as a bargaining problem, 2026. URL <https://arxiv.org/abs/2602.04155>.
- Wassily Hoeffding. Probability inequalities for sums of bounded random variables. *Journal of the American Statistical Association*, 58(301):13–30, 1963.
- Ehud Kalai and Meir Smorodinsky. Other solutions to nash’s bargaining problem. *Econometrica*, 43(3):513–518, 1975.
- Rémi Munos, Michal Valko, Daniele Calandriello, Mohammad Gheshlaghi Azar, Mark Rowland, Daniel Guo, Yunhao Tang, Matthieu Geist, Thomas Mesnard, Andrea Michi, et al. Nash learning from human feedback. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, 2024.
- John F Nash. The bargaining problem. *Econometrica*, 18(2):155–162, 1950.
- Aviv Navon, Aviv Shamsian, Idan Achituve, Haggai Maron, Kenji Kawaguchi, Gal Chechik, and Ethan Fetaya. Multi-task learning as a bargaining game. In *Proceedings of the 39th International Conference on Machine Learning*, volume 162 of *Proceedings of Machine Learning Research*, 2022.
- Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, John Schulman, Jacob Hilton, Fraser Kelton, Luke Miller, Maddie Simens, Amanda Askell, Peter Welinder, Paul Christiano, Jan Leike, and Ryan Lowe. Training language models to follow instructions with human feedback. In *Advances in Neural Information Processing Systems*, volume 35, pp. 27730–27744, 2022.
- Rafael Rafailov, Archit Sharma, Eric Mitchell, Stefano Ermon, Christopher D. Manning, and Chelsea Finn. Direct preference optimization: Your language model is secretly a reward model. In *Advances in Neural Information Processing Systems*, volume 36, 2023.
- Alexandre Ramé, Guillaume Couairon, Corentin Dancette, Jean-Baptiste Gaya, Mustafa Shukor, Laure Soulier, and Matthieu Cord. Rewarded soups: Towards pareto-optimal alignment by interpolating weights fine-tuned on diverse rewards. In *Advances in Neural Information Processing Systems*, volume 36, 2023.
- Corby Rosset, Ching-An Cheng, Arindam Mitra, Michael Santacroce, Ahmed Awadallah, and Tengyang Xie. Direct nash optimization: Teaching language models to self-improve with general preferences. *arXiv preprint arXiv:2404.03715*, 2024.
- Seongho Son, William Bankes, Sangwoong Yoon, Shyam Sundhar Ramesh, Xiaohang Tang, and Ilija Bogunovic. Robust multi-objective controlled decoding of large language models. In *International Conference on Learning Representations*, 2026.
- Benjamin Warner, Antoine Chaffin, Benjamin Clavié, et al. Smarter, better, faster, longer: A modern bidirectional encoder for fast, memory efficient, and long context finetuning and inference, 2024. URL <https://arxiv.org/abs/2412.13663>.
- Yue Wu, Zhiqing Sun, Huizhuo Yuan, Kaixuan Ji, Yiming Yang, and Quanquan Gu. Self-play preference optimization for language model alignment. In *International Conference on Learning Representations*, 2025.
- Jiahao Zhang, Lujing Zhang, Keltin Grimes, Zhuohao Yu, Gokul Swamy, and Zhiwei Steven Wu. Back to blackwell: Closing the loop on intransitivity in multi-objective preference fine-tuning. *arXiv preprint arXiv:2602.19041*, 2026.

Yuheng Zhang, Dian Yu, Baolin Peng, Linfeng Song, Ye Tian, Mingyue Huo, Nan Jiang, Haitao Mi, and Dong Yu. Iterative nash policy optimization: Aligning LLMs with general preferences via no-regret learning. In *The Thirteenth International Conference on Learning Representations*, 2025. URL <https://openreview.net/forum?id=Pujt3ADZgI>.

Huiying Zhong, Tianwei Gao, Zhiwei Steven Wu, Linjun Zhang, Weijie J. Su, and Zhun Deng. Provable pluralistic alignment: Multi-party rlhf under offline human feedback. *arXiv preprint arXiv:2403.05006v2*, 2026. URL <https://arxiv.org/abs/2403.05006v2>.

Yifan Zhong, Chengdong Ma, Xiaoyuan Zhang, Ziran Yang, Haojun Chen, Qingfu Zhang, Siyuan Qi, and Yaodong Yang. Panacea: Pareto alignment via preference adaptation for LLMs. In *Advances in Neural Information Processing Systems*, 2024.

Runlong Zhou, Maryam Fazel, and Simon S. Du. Extragradients preference optimization (EGPO): Beyond last-iterate convergence for nash learning from human feedback. In *Conference on Language Modeling*, 2025.

Zhanhui Zhou, Jie Liu, Jing Shao, Xiangyu Yue, Chao Yang, Wanli Ouyang, and Yu Qiao. Beyond one-preference-fits-all alignment: Multi-objective direct preference optimization. In *Findings of the Association for Computational Linguistics: ACL 2024*, 2024.

## A PROOFS

### A.1 PROOF OF PROPOSITION 1

For a fixed comparator  $\nu$ , define

$$\mathcal{J}_{k,\nu}(\pi) := J_k(\pi, \nu) + \beta_k \mathbb{E}_{x \sim \mathcal{D}} \text{KL}(\nu(\cdot | x) \| \pi_{\text{ref}}(\cdot | x)).$$

The first term is linear in  $\pi$  and the second does not depend on  $\pi$ , so  $\mathcal{J}_{k,\nu}$  is affine. The pointwise minimum of affine functions is concave; hence  $V_{k,\beta_k}$  is concave. Since  $|\Delta_k| \leq 1/2$ ,

$$|V_{k,\beta_k}(\pi) - V_{k,\beta_k}(\pi')| \leq \frac{1}{2} \mathbb{E}_{x \sim \mathcal{D}} \|\pi(\cdot | x) - \pi'(\cdot | x)\|_1,$$

which also proves continuity.

For  $\beta_k > 0$ , the inner minimization separates across prompts. Fix  $x$  and abbreviate  $m(z) = m_k^\pi(z | x)$ . The prompt-wise problem is

$$\min_{\nu \in \Delta(\mathcal{Y}_x)} \sum_z \nu(z) m(z) + \beta_k \sum_z \nu(z) \log \frac{\nu(z)}{\pi_{\text{ref}}(z | x)}.$$

The first-order condition under  $\sum_z \nu(z) = 1$  gives

$$\nu(z) \propto \pi_{\text{ref}}(z | x) \exp\{-m(z)/\beta_k\},$$

which yields (8); substituting it back yields (9). The KL term makes the minimizer unique. Danskin’s theorem therefore gives, for every feasible direction  $h$  satisfying  $\sum_y h(y | x) = 0$ ,

$$DV_{k,\beta_k}(\pi)[h] = \mathbb{E}_{x \sim \mathcal{D}} \sum_y h(y | x) \mathbb{E}_{z \sim \nu_{k,\pi}^*(\cdot | x)} \Delta_k(y, z | x).$$

Thus the coordinate derivative is

$$[\nabla_\pi V_{k,\beta_k}(\pi)]_{x,y} = \mathcal{D}(x) Q_k^\pi(y | x),$$

proving (10).

### A.2 CONNECTION TO THE FIXED-REFERENCE

This subsection gives a finite-action illustration at the unregularized endpoint  $\beta = 0$ . It is not used in Assumption 1, the full-support analysis, or the finite-temperature experiments.

Define the fixed-reference improvement

$$a_k(\pi) = J_k(\pi, \pi_{\text{ref}}) = P_k(\pi \succ \pi_{\text{ref}}) - \frac{1}{2}. \quad (28)$$

Skew-symmetry gives  $a_k(\pi_{\text{ref}}) = 0$ .

**Full cyclic example.** For one prompt, let  $\pi_{\text{ref}} = \delta_r$ , where  $\delta_y$  is the deterministic policy that always outputs  $y$ . Suppose  $A$ ,  $B$ , and  $C$  each beat  $r$  with probability 0.75 and form the deterministic cycle

$$A \succ B, \quad B \succ C, \quad C \succ A.$$

Both

$$\pi_{\text{pure}} = \delta_A \quad \text{and} \quad \pi_{\text{mix}} = \frac{\delta_A + \delta_B + \delta_C}{3}$$

have the same fixed-reference improvement,

$$a(\pi_{\text{pure}}) = a(\pi_{\text{mix}}) = 0.75 - \frac{1}{2} = 0.25.$$

Their worst-opponent values are nevertheless different. For the pure policy, response  $C$  defeats  $A$  deterministically, so

$$V_0(\pi_{\text{pure}}) = \min_{\nu} J(\pi_{\text{pure}}, \nu) = J(\delta_A, \delta_C) = -\frac{1}{2}.$$

For the uniform mixture, consider first the pure opponent  $\delta_A$ . The mixture outputs  $A$ ,  $B$ , and  $C$  with equal probability. Against  $A$ , these responses respectively tie, lose, and win, giving centered payoffs 0,  $-1/2$ , and  $+1/2$ . Hence

$$J(\pi_{\text{mix}}, \delta_A) = \frac{1}{3} \left( 0 - \frac{1}{2} + \frac{1}{2} \right) = 0.$$

By symmetry,

$$J(\pi_{\text{mix}}, \delta_B) = J(\pi_{\text{mix}}, \delta_C) = 0.$$

Against the reference,

$$J(\pi_{\text{mix}}, \delta_r) = 0.25,$$

because every response in the mixture beats  $r$  with probability 0.75. Since  $J(\pi_{\text{mix}}, \nu)$  is linear in  $\nu$ , no mixed opponent can achieve a value below the minimum over these pure opponents. Therefore

$$V_0(\pi_{\text{mix}}) = \min_{\nu} J(\pi_{\text{mix}}, \nu) = \min\{0, 0, 0, 0.25\} = 0.$$

Finally, the reference itself has

$$V_0(\pi_{\text{ref}}) = \min_{\nu} J(\delta_r, \nu) = -0.25,$$

since  $r$  loses to each of  $A$ ,  $B$ , and  $C$  with probability 0.75. Thus

|                     | $a(\pi)$ | $V_0(\pi)$ | $V_0(\pi) - V_0(\pi_{\text{ref}})$ |
|---------------------|----------|------------|------------------------------------|
| $\pi_{\text{pure}}$ | 0.25     | -0.50      | -0.25                              |
| $\pi_{\text{mix}}$  | 0.25     | 0          | 0.25                               |

The fixed reference therefore assigns the two policies the same score, even though the pure policy is fully exploitable by one response whereas the balanced mixture guarantees an average tie against every response in the cycle.

**Game-value improvement need not imply a direct reference win.** For one prompt, consider

$$A = \begin{pmatrix} 0 & .5 & -.5 \\ -.5 & 0 & .5 \\ .5 & -.5 & 0 \end{pmatrix}, \quad \mu = (.8, .1, .1), \quad p = (.25, .43, .32), \quad \beta = .25.$$

Both policies have full support. Direct evaluation gives  $V_{\beta}(\mu) = -.0517901$  and  $V_{\beta}(p) = -.0417293$ , hence  $s(p) = .0100608 > 0$ . Yet  $p^{\top} A \mu = -.0385$ , so the direct reference win probability is .4615. Robustness-adjusted value and direct wins answer different questions even under the same judge.

### A.3 POLICY-INDUCED NASH BARGAINING AND CLASSICAL CONSEQUENCES

We prove Proposition 2. Since  $\Pi$  is compact and each  $V_{k,\beta_k}$  is continuous,  $V_\beta(\pi)$  is bounded. The definition of  $\mathcal{G}_\beta$  then implies boundedness. Closedness follows by taking any convergent sequence  $v_n \in \mathcal{G}_\beta$ , choosing policies  $\pi_n$  with  $d \leq v_n \leq V_\beta(\pi_n)$ , extracting a convergent policy subsequence, and applying continuity. Downward closure within  $\{v : v \geq d\}$  is immediate from the definition.

For convexity, take  $v^1, v^2 \in \mathcal{G}_\beta$  with witnesses  $\pi_1, \pi_2$ . Objective-wise concavity from Proposition 1 gives, for  $\alpha \in [0, 1]$ ,

$$V_\beta(\alpha\pi_1 + (1 - \alpha)\pi_2) \geq \alpha V_\beta(\pi_1) + (1 - \alpha)V_\beta(\pi_2)$$

componentwise, so  $\alpha v^1 + (1 - \alpha)v^2 \in \mathcal{G}_\beta$ . Assumption 1 provides a point strictly above  $d$ , proving nondegeneracy.

For any  $v \in \mathcal{G}_\beta$  with  $v > d$ , some policy  $\pi$  satisfies  $v \leq V_\beta(\pi)$ . Since the Nash objective is coordinatewise increasing,

$$\sum_k \log(v_k - d_k) \leq \sum_k \log(V_{k,\beta_k}(\pi) - d_k) = F(\pi).$$

Conversely, every  $\pi \in \Pi_+$  induces a feasible value vector  $V_\beta(\pi)$ . Hence (12) realizes the classical Nash solution over  $(\mathcal{G}_\beta, d)$ .

The remaining statements are standard Nash-bargaining consequences, but we record the short arguments for completeness. Since the log objective tends to  $-\infty$  when any surplus approaches zero, every optimizer lies strictly above  $d$ . Coordinatewise monotonicity implies Pareto optimality. Strict concavity of  $\sum_k \log s_k$  over value vectors implies that all optimal policies induce the same surplus, hence the same game-value vector, although the policy itself can be nonunique.

Finally, under independent positive affine transformations

$$V'_k(\pi) = c_k V_k(\pi) + b_k, \quad d'_k = c_k d_k + b_k, \quad c_k > 0,$$

the transformed surplus is  $s'_k(\pi) = c_k s_k(\pi)$  and

$$\sum_k \log s'_k(\pi) = \sum_k \log s_k(\pi) + \sum_k \log c_k.$$

The added term is independent of  $\pi$ , so the optimizer set is unchanged.

Table 3: Guarantees of the stated population rules under joint improvement. A “no” means the property need not hold for every optimizer without an additional tie-break. This is not an exhaustive comparison of bargaining solutions.

| rule                        | strict IR | Pareto optimal | unique value | affine invariant |
|-----------------------------|-----------|----------------|--------------|------------------|
| positive-weight utilitarian | no        | yes            | no           | no               |
| absolute maxmin             | no        | no             | no           | no               |
| surplus maxmin              | yes       | no             | no           | no               |
| Nash bargaining             | yes       | yes            | yes          | yes              |

**Scope of invariance and alternatives.** Affine invariance transforms completed utilities and disagreement values together. At the primitive level,  $A'_k = c_k A_k$  requires  $\beta'_k = c_k \beta_k$  to preserve the opponent and obtain  $V'_k = c_k V_k$ . With fixed temperature instead,  $V_{cA,\beta} = cV_{A,\beta/c}$ . If scaled margins leave  $[-1/2, 1/2]$ , the transformation changes payoff units rather than defining a new probability oracle. Arbitrary judge recalibration is not covered.

Exact dual multipliers transform as  $\lambda'_k = \lambda_k/c_k$ , leaving  $\lambda'_k Q'_k = \lambda_k Q_k$  and the pair target unchanged. Matching projected-gradient trajectories further requires scaling initial multipliers and coordinate bounds by  $1/c_k$ , any diagnostic surplus thresholds by  $c_k$ , and coordinate step sizes by  $1/c_k^2$ . Target signs alone do not verify these stronger identities. We do not floor surpluses to make infeasible targets positive.

A Pareto-refining tie-break can remove the dominated-optimum failure of raw maxmin. Relative-improvement bargaining provides another scale-invariant choice but requires objective-wise ideal values and, in general, an appropriate refinement (Kalai & Smorodinsky, 1975; Han et al., 2026). These alternatives do not make Nash uniquely admissible. Here the bargaining parties are objective rubrics: duplicating a rubric duplicates its log-surplus term. The fixed reference also plays two distinct roles, disagreement policy and opponent anchor; changing both changes the utilities as well as disagreement.

**Rubric and reference sensitivity (pending).** This check is declared and unmeasured. It prespecifies a duplication or merge of an overlapping rubric and reports its effect on the selected game values and on policy outcomes; the first half needs only a re-solve on the frozen teacher values, the second needs a training run at the matched budget. Separately vary the disagreement policy while holding the opponent anchor fixed, and vary the opponent anchor while holding the disagreement policy fixed. Recompute the disagreement values under each corresponding game definition; a fixed disagreement policy need not have fixed game values when the anchor changes. These checks assess modeling choices that affine invariance does not remove.

#### A.4 COUNTEREXAMPLES FOR ALTERNATIVE AGGREGATION RULES

The positive guarantees in Table 3 follow directly from the corresponding objectives. Because  $w_k > 0$ , the positive-weight utilitarian objective  $\sum_k w_k V_k(\pi)$  is strictly increasing in every game-value coordinate, so none of its optimizers can be Pareto dominated. Under Assumption 1, surplus maxmin also attains a strictly positive minimum surplus because a jointly improving policy exists. The examples below establish the guarantees that can fail for the alternative rules.

All examples use three responses, so a policy is a probability vector  $\pi = (\pi_1, \pi_2, \pi_3) \in \Delta_3$ , and use the uniform reference  $\pi_{\text{ref}} = (1/3, 1/3, 1/3)$ . For objective  $k$ ,  $\Delta_k$  is the skew-symmetric pairwise payoff matrix:  $(\Delta_k)_{ij}$  is the centered preference margin of response  $i$  against response  $j$ . At  $\beta_k = 0$ ,

$$V_k(\pi) = \min_j (\pi^\top \Delta_k)_j, \quad d_k = V_k(\pi_{\text{ref}}), \quad s_k(\pi) = V_k(\pi) - d_k.$$

Thus  $\pi^\top \Delta_k$  contains the expected payoff of  $\pi$  against each pure opponent, and  $V_k(\pi)$  is its smallest coordinate.

Because each  $V_k$  is the minimum of linear functions, the optimizers below are obtained from small linear programs. For example, equal-weight utilitarian optimization can be written as

$$\max_{\pi, t_1, t_2} t_1 + t_2 \quad \text{s.t.} \quad t_k \leq (\pi^\top \Delta_k)_j \quad \forall k, j, \quad \pi \in \Delta_3.$$

Absolute maxmin uses one variable  $t$  with  $t \leq (\pi^\top \Delta_k)_j$ , while surplus maxmin uses  $t + d_k \leq (\pi^\top \Delta_k)_j$  for every  $k, j$ .

**Utilitarian failure of strict improvement.** Let

$$\Delta_1 = \begin{pmatrix} 0 & -1/2 & -1/2 \\ 1/2 & 0 & -1/2 \\ 1/2 & 1/2 & 0 \end{pmatrix}, \quad \Delta_2 = \begin{pmatrix} 0 & -1/2 & 1/4 \\ 1/2 & 0 & -1/2 \\ -1/4 & 1/2 & 0 \end{pmatrix}.$$

For the uniform reference,

$$\pi_{\text{ref}}^\top \Delta_1 = (1/3, 0, -1/3), \quad \pi_{\text{ref}}^\top \Delta_2 = (1/12, 0, -1/12),$$

and hence

$$d = (-1/3, -1/12).$$

First, joint improvement is feasible. The policy

$$\bar{\pi} = (1/4, 1/4, 1/2)$$

is only a feasibility witness; it gives

$$\bar{\pi}^\top \Delta_1 = (3/8, 1/8, -1/4), \quad \bar{\pi}^\top \Delta_2 = (0, 1/8, -1/16).$$

Therefore

$$V(\bar{\pi}) = (-1/4, -1/16), \quad s(\bar{\pi}) = V(\bar{\pi}) - d = (1/12, 1/48) > 0.$$

In contrast, solving the equal-weight utilitarian problem  $\max_{\pi} V_1(\pi) + V_2(\pi)$  gives

$$\pi_U = (0, 1/5, 4/5).$$

For this policy,

$$\pi_U^\top \Delta_1 = (1/2, 2/5, -1/10), \quad \pi_U^\top \Delta_2 = (-1/10, 2/5, -1/10),$$

so

$$V(\pi_U) = (-1/10, -1/10), \quad s(\pi_U) = (7/30, -1/60).$$

Thus joint improvement is possible, yet a utilitarian optimum falls below the fallback on objective 2.

**Absolute-maxmin failure of strict improvement.** Let

$$\Delta_1 = \begin{pmatrix} 0 & -1/2 & 1/4 \\ 1/2 & 0 & 1/2 \\ -1/4 & -1/2 & 0 \end{pmatrix}, \quad \Delta_2 = \begin{pmatrix} 0 & 1/4 & -1/2 \\ -1/4 & 0 & 1/2 \\ 1/2 & -1/2 & 0 \end{pmatrix}.$$

Again,

$$d = (-1/3, -1/12).$$

Joint improvement is feasible, as witnessed by

$$\bar{\pi} = (1/3, 7/15, 1/5),$$

for which

$$V(\bar{\pi}) = (-4/15, -1/60)$$

and therefore

$$s(\bar{\pi}) = (1/15, 1/15) > 0.$$

Absolute maxmin instead solves

$$\max_{\pi} \min\{V_1(\pi), V_2(\pi)\},$$

whose linear-program solution is

$$\pi_{AM} = (0, 4/5, 1/5).$$

For this policy,

$$\pi_{AM}^\top \Delta_1 = (7/20, -1/10, 2/5), \quad \pi_{AM}^\top \Delta_2 = (-1/10, -1/10, 2/5),$$

so

$$V(\pi_{AM}) = (-1/10, -1/10)$$

and

$$s(\pi_{AM}) = (7/30, -1/60).$$

Thus absolute maxmin maximizes the smallest *absolute* game value while allowing objective 2 to fall below its own fallback.

**A Pareto-dominated surplus-maxmin optimum.** Let

$$\Delta_1 = \begin{pmatrix} 0 & -1/2 & 1/2 \\ 1/2 & 0 & 1/2 \\ -1/2 & -1/2 & 0 \end{pmatrix}, \quad \Delta_2 = \begin{pmatrix} 0 & 0 & 1/2 \\ 0 & 0 & 1/2 \\ -1/2 & -1/2 & 0 \end{pmatrix}.$$

The uniform reference gives

$$d = (-1/3, -1/6).$$

For objective 2, no policy can achieve  $V_2(\pi) > 0$ , while every mixture supported on the first two responses achieves  $V_2(\pi) = 0$ . Consequently,

$$s_2(\pi) \leq 1/6$$

for every policy, so the optimal surplus-maxmin value cannot exceed  $1/6$ .

The two policies

$$\pi_c = (1/3, 2/3, 0), \quad \pi_e = (0, 1, 0)$$

satisfy

$$V(\pi_c) = (-1/6, 0), \quad s(\pi_c) = (1/6, 1/6),$$

and

$$V(\pi_e) = (0, 0), \quad s(\pi_e) = (1/3, 1/6).$$

Both therefore maximize

$$\min_k s_k(\pi)$$

at  $1/6$ . However,

$$s(\pi_e) > s(\pi_c)$$

in the first coordinate and is equal in the second, so  $\pi_e$  Pareto dominates  $\pi_c$ . Thus the surplus-maxmin optimizer set can contain Pareto-dominated policies.

## A.5 DUAL DERIVATION, DIRECT SOLVE, AND NEURAL PROJECTION

For multipliers  $\lambda_k \geq 0$ , the Lagrangian of (16) is

$$\mathcal{L}_t(\pi, u, \lambda) = \sum_k \log u_k + \sum_k \lambda_k (s_k(\pi) - u_k) - \eta_t^{-1} D(\pi \| \pi_t).$$

For fixed  $\lambda_k > 0$ , maximizing  $\log u_k - \lambda_k u_k$  gives  $u_k = 1/\lambda_k$  and value  $-\log \lambda_k - 1$ ; if  $\lambda_k = 0$ , the maximum is unbounded. Substitution yields (17), and the remaining maximization over  $\pi$  is (18). Uniqueness of that maximizer and Danskin’s theorem yield (19).

Assumption 1 provides strict feasibility: choose  $\bar{\pi}$  with  $s_k(\bar{\pi}) > 0$  and  $0 < \bar{u}_k < s_k(\bar{\pi})$ . Slater’s condition therefore gives strong duality. Stationarity in  $u_k$  gives

$$\lambda_{t+1,k} = \frac{1}{u_{t+1,k}} > 0,$$

and complementary slackness then forces  $u_{t+1,k} = s_k(\pi_{t+1})$ . This proves Proposition 3.

**Dual step used as a fallback.** Stationarity  $\hat{s}_k(p_t^*(\lambda)) = 1/\lambda_k$  can also be approached by projected subgradient *descent* on the dual — the update below subtracts the dual gradient, and the dual of the concave inner problem is minimized over  $\lambda$  —

$$\lambda^{(m+1)} = \Pi_\Lambda \left[ \lambda^{(m)} - \gamma_m \left( \hat{s} \left( p_t^* \left( \lambda^{(m)} \right) \right) - \left( \lambda^{(m)} \right)^{-1} \right) \right], \quad (29)$$

where  $\gamma_m > 0$  is the dual step size,  $\Pi_\Lambda$  is Euclidean projection onto  $\Lambda = [\lambda_{\min}, \lambda_{\max}]^K$  with  $0 < \lambda_{\min} < \lambda_{\max}$ , and the inverse is elementwise. On the pools reported here this iteration stalls at a projected KKT residual near  $10^{-2}$  after 3000 steps, whereas the root solve reaches  $10^{-9}$  in tens of inner solves; no reported result uses (29).

**What the projected certificate certifies.** The projected residual concerns the box-constrained dual. Original stationarity additionally requires  $\|\hat{s} - \lambda^{-1}\|_\infty$  to meet tolerance; active bounds cannot be ignored merely because projection removes the residual. We record each bound hit and classify a solve with unresolved unprojected stationarity as uncertified. Interior numerical solutions with small residuals support the finite problem only. The supplied SafeRLHF report gives  $\Lambda = [10^{-3}, 10^3]^2$ ,  $\lambda = (6.33, 6.19)$ , and independently recomputed unprojected residual  $6.7 \times 10^{-11}$ . These are teacher diagnostics, not neural-policy certificates.

**Finite-pool optimality and neural projection.** Fix  $\lambda$  and let  $\pi_t^*(\lambda)$  solve (18). First-order optimality gives, for every prompt and response pair in the support,

$$h_t(\pi_t^*(\lambda); x, y, y') = \eta_t \sum_k \lambda_k \left( Q_k^{\pi_t^*(\lambda)}(y | x) - Q_k^{\pi_t^*(\lambda)}(y' | x) \right). \quad (30)$$

The direct finite-pool program (23) is the sampled analogue of this strictly concave problem. Its normalized optimizer determines all finite-pool log-probability differences, so (24) is identified without computing a language-model normalizer. Equation (27) is the corresponding finite-pool KKT identity and is checked numerically before an artifact is accepted. If the positive-probability pair graph connects all sampled occurrences, zero regression error implies

$$\tilde{p}_{\theta,x}(i) = \frac{p_{t,x}(i) \pi_{\theta}(y_{x,i} | x) / \pi_t(y_{x,i} | x)}{\sum_j p_{t,x}(j) \pi_{\theta}(y_{x,j} | x) / \pi_t(y_{x,j} | x)} = p_{t,x}^*(i).$$

This identifies the importance-reweighted occurrence measure only. For example, under a uniform three-response anchor, fitting only pair (1, 2) with target zero permits every neural distribution  $(a, a, 1 - 2a)$ ,  $0 < a < 1/2$ . Unseen-response mass remains unconstrained. Resampling a new learner pool at each outer stage makes the empirical base measure track  $\pi_t$ ; reusing a pool from another proposal instead requires explicit importance weights. The reference opponent anchor remains  $\pi_{\text{ref}}$ .

**Why an optimal fixed point need not be reached.** Use the cyclic matrix  $A$  above, reference and proximal center  $\mu = (.34, .33, .33)$ ,  $\beta = .25$ , and fixed  $\eta\lambda = 4$ . Undamped best-response alternation applies

$$T(p) = \text{softmax}(\log \mu + 4A \text{softmax}(\log \mu + Ap/\beta)).$$

Its fixed point is  $p^* = (.33618258, .32907171, .33474571)$ , but both tangent-space Jacobian eigenvalues are approximately  $-1.333163$ . Perturbations therefore alternate and expand locally. This is not excluded by improvement:  $V(p^*) - V(\mu) \simeq 4.7194 \times 10^{-5} > 0$ . In the supplied deterministic check, 100 iterations from  $\mu$  leave  $\|p - p^*\|_2 \simeq .27986$ ; a direct constrained solve of the same fixed-multiplier problem reaches inner stationarity below  $10^{-9}$ . These are finite-game calculations, not LLM results or a complete Nash dual solve.

For general fixed  $\lambda$ , let  $f_x(p; \lambda)$  denote the objective in (23) and let  $\nu_k(p) = \text{softmax}(\log \mu_x - \hat{A}_{k,x}^\top p / \beta_k)$ . On the simplex interior,

$$\begin{aligned} \nabla f_x(p; \lambda) &= \sum_k \lambda_k \hat{A}_{k,x} \nu_k(p) - \eta_t^{-1} (\log(p/p_{t,x}) + \mathbf{1}), \\ \nabla^2 f_x(p; \lambda) &= - \sum_k \frac{\lambda_k}{\beta_k} \hat{A}_{k,x} (\text{diag}(\nu_k) - \nu_k \nu_k^\top) \hat{A}_{k,x}^\top - \eta_t^{-1} \text{diag}(1/p) \prec 0. \end{aligned}$$

Thus the finite target is unique, and direct optimization includes opponent feedback through this curvature. Strict concavity does not certify an arbitrary numerical routine: simplex feasibility and inner/dual residuals still require measurement. The empirical tensor itself remains an approximation to population preferences; exponentiation, normalization, and log-sum-exp do not preserve unbiasedness of noisy margins.

**Finite-pool feasibility and failure reporting.** When diagnosing an uncertified solve, use the finite-pool program

$$\hat{\tau}^* = \max_{\{p_x\}, \tau} \tau \quad \text{s.t.} \quad \mathbb{E}_{x \sim \hat{\mathcal{D}}} \hat{V}_{k,x}(p_x) - \hat{d}_k \geq \tau \quad (\forall k), \quad p_x \in \Delta^{n_x}.$$

This is convex optimization with concave-superlevel constraints, not generally a linear program at  $\beta_k > 0$ . A verified positive feasible margin witnesses strict finite-pool improvement. Only a valid nonpositive upper bound certifies its absence; failure to find a witness is unresolved. A tolerance-bounded interval near zero means no resolved positive margin, not exact infeasibility. Independent learner and reference banks need not contain a categorical representation of the fallback, so finite-pool infeasibility is not population infeasibility.

For an analytic infeasible example, take two actions,  $\mu = (1/2, 1/2)$ ,  $A_1 = \begin{pmatrix} 0 & 1/2 \\ -1/2 & 0 \end{pmatrix}$  and  $A_2 = -A_1$ . For  $p = (a, 1 - a)$  and any common  $\beta > 0$ ,  $s_1(p) = (a - 1/2)/2 = -s_2(p)$ , hence  $\tau^* = 0$ . No clipping or inverse-surplus floor creates a strictly feasible Nash target.

The execution contract records feasible, certified-infeasible, numerical-unresolved, and neural-rejected outcomes separately. Uncertified solves retain the previous policy. This audit introduces no separate neural warm-start stage. Post-fit acceptance uses prespecified development checks, including empirical surplus signs, target fit, and numerical validity; reused monitoring prompts are validation data. False acceptance/rejection rates near the boundary remain an unmeasured audit, and generic bounded-mean confidence formulas do not automatically cover biased game-value estimates.

## A.6 LAST-ITERATE CONVERGENCE

**Theorem 1** (Last-iterate convergence). *Under Assumption 1, let  $\pi^*$  maximize (12), set  $\pi_0 = \pi_{\text{ref}}$ , and solve (16) exactly at every outer iteration with arbitrary  $\eta_t > 0$ . Then  $F(\pi_{t+1}) \geq F(\pi_t)$  for  $t \geq 1$ , and for every  $T \geq 1$ ,*

$$F(\pi^*) - F(\pi_T) \leq \frac{D(\pi^* \| \pi_{\text{ref}})}{\sum_{t=0}^{T-1} \eta_t}. \quad (31)$$

Work throughout in the  $\mathcal{D}$ -weighted geometry: pair functions of  $(x, y)$  by

$$\langle f, g \rangle_{\mathcal{D}} = \mathbb{E}_{x \sim \mathcal{D}} \sum_y f(x, y) g(x, y),$$

and let  $h$  be the *prompt-wise* negative entropy  $h(\pi)(x) = \sum_y \pi(y | x) \log \pi(y | x)$ , whose gradient at  $(x, y)$  is  $1 + \log \pi(y | x)$ . Its Bregman divergence in that geometry is exactly  $D(\pi \| \pi_t) = \mathbb{E}_{x \sim \mathcal{D}} \text{KL}(\pi(\cdot | x) \| \pi_t(\cdot | x))$ . The prompt expectation enters once, through the pairing, and is not also carried inside  $h$ . Because the objective is increasing in  $u$ , the policy part of (16) is equivalently

$$\pi_{t+1} \in \arg \max_{\pi \in \Pi_+} \left\{ F(\pi) - \frac{1}{\eta_t} D(\pi \| \pi_t) \right\}. \quad (32)$$

The feasible set is nonempty by Assumption 1, and  $F(\pi) \rightarrow -\infty$  as any surplus approaches zero. Thus the maximizer exists in  $\Pi_+$  even when the center  $\pi_t$  is the reference  $\pi_{\text{ref}} \notin \Pi_+$ . Because  $\pi_t$  has full support and the remaining objective has finite directional derivatives at an interior surplus point, the inward derivative of  $-D(\pi \| \pi_t)$  excludes zero-probability coordinates. Induction therefore gives a full-support, positive-surplus  $\pi_{t+1}$ .

Let  $g_{t+1}$  be a supergradient of  $F$  at  $\pi_{t+1}$ . First-order optimality of (32) and concavity of  $F$  imply, for every  $\pi \in \Pi_+$ ,

$$F(\pi) - F(\pi_{t+1}) \leq \langle g_{t+1}, \pi - \pi_{t+1} \rangle_{\mathcal{D}} \quad (33)$$

$$\leq \frac{1}{\eta_t} \langle \nabla h(\pi_{t+1}) - \nabla h(\pi_t), \pi - \pi_{t+1} \rangle_{\mathcal{D}} \quad (34)$$

$$= \frac{1}{\eta_t} [D(\pi \| \pi_t) - D(\pi \| \pi_{t+1}) - D(\pi_{t+1} \| \pi_t)]. \quad (35)$$

For  $t \geq 1$ ,  $\pi_t$  is feasible in the next subproblem. Comparing the optimum with  $\pi_t$  gives

$$F(\pi_{t+1}) - \frac{1}{\eta_t} D(\pi_{t+1} \| \pi_t) \geq F(\pi_t), \quad (36)$$

so the objective values are nondecreasing from  $\pi_1$  onward.

Set  $\pi = \pi^*$  in (35), multiply by  $\eta_t$ , and sum from  $t = 0$  to  $T - 1$ :

$$\sum_{t=0}^{T-1} \eta_t (F(\pi^*) - F(\pi_{t+1})) \leq D(\pi^* \| \pi_{\text{ref}}). \quad (37)$$

Because the gaps are nonincreasing, the left-hand side is at least  $(\sum_t \eta_t)(F(\pi^*) - F(\pi_T))$ , which proves (31).




Figure 2: UF-4 method means from Table 1 on 1,236 common prompts; bars show one sample SD over policy seeds, not prompt intervals. PROSPER has two seeds and other trained rows have three. Projections do not establish full-dimensional Pareto dominance; the DPO weight sweep is incomplete.

**Corollary 1** (Convergence of the surplus profile). *Under the conditions of Theorem 1, let  $s^* = s(\pi^*)$  denote the unique surplus vector induced by a Nash-optimal policy. Then*

$$\|s(\pi_T) - s^*\|_2^2 \leq \frac{2D(\pi^* \|\pi_{\text{ref}})}{\sum_{t=0}^{T-1} \eta_t}. \quad (38)$$

Hence a constant proximal step gives  $\|s(\pi_T) - s^*\|_2 = O(T^{-1/2})$ .

*Proof.* This is a stronger consequence of Nash-welfare convergence rather than a separate requirement for the NBPO update. For every objective,  $V_k(\pi) \in [-1/2, 1/2]$ , and  $d_k = V_k(\pi_{\text{ref}}) \leq 0$  because  $\nu = \pi_{\text{ref}}$  is feasible and  $J_k(\pi_{\text{ref}}, \pi_{\text{ref}}) = 0$ . Hence every positive surplus is at most one. On  $(0, 1]^K$ , the function  $G(s) = \sum_k \log s_k$  is 1-strongly concave in Euclidean norm. Proposition 2 implies that all Nash-optimal policies induce the same game-value vector, and therefore the same surplus vector  $s^*$ . First-order optimality over the convex comprehensive surplus set induced by (14) gives

$$G(s^*) - G(s) \geq \frac{1}{2} \|s - s^*\|_2^2. \quad (39)$$

Applying this inequality to  $s = s(\pi_T)$  and using (31) proves (38).  $\square$

## B ADDITIONAL POLICY AND REALIZATION RESULTS

The following measured results and diagnostics complement the main comparison; they retain their original evaluation contracts and reported numerical entries.

### B.1 CAPABILITY, REALIZATION, AND SUPPORTING RESULTS

**Capability retention.** Table 4 reports UF checkpoint outcomes. NBPO’s IFEval strict-prompt accuracy falls from the base’s 0.7542 to 0.7264; one-seed DPO reaches 0.6359 despite higher judged IF. MOPO reaches 0.7579. This discrepancy requires constraint-type and length analysis. GSM8K and the appendix’s ARC-C, HellaSwag and MMLU change relatively little, without establishing equivalence.

HarmBench harmful completion rises from 0.2812 to 0.3906–0.4333 across matched mechanisms and 0.6281 for one-seed DPO. XSTest refusal falls on both safe and unsafe prompts. These observations are consistent with increased compliance; harmlessness is not among the four UF objectives. Non-refusal alone is not harmful compliance, and direct-request HarmBench is not an attack-search evaluation.

DPO leads the reported local AlpacaEval-2, Arena-Hard-v2 and MT-Bench proxies while its safety and IFEval outcomes deteriorate; these proxies are not official leaderboard scores. MOPO retains

Table 4: UF-4 capability and safety; sample SD is below multi-seed rows. IFEval: strict-prompt accuracy; GSM8K: exact match; HB: harmful completion; XS: safe-prompt full refusal. AE2/AH2/MT are local proxies, not official leaderboard scores. DPO has one capability seed. Coverage and protocols: Appendix G.1.

| Method                      | Seeds | IFEval | GSM8K  | HB↓    | XS↓    | AE2†   | AH2†   | MT†  |
|-----------------------------|-------|--------|--------|--------|--------|--------|--------|------|
| Base                        | 1     | 0.7542 | 0.8666 | 0.2812 | 0.0683 | —      | —      | 7.53 |
| NBPO                        | 3     | 0.7264 | 0.8633 | 0.3906 | 0.0414 | 0.5617 | 0.5143 | 7.59 |
| <i>sample SD over seeds</i> |       | 0.0098 | 0.0031 | 0.0094 | 0.0023 | 0.0083 | 0.0206 | 0.09 |
| Fixed-reference Nash        | 3     | 0.7289 | 0.8615 | 0.3990 | 0.0360 | 0.5578 | 0.5243 | 7.60 |
| <i>sample SD over seeds</i> |       | 0.0059 | 0.0038 | 0.0072 | 0.0001 | 0.0138 | 0.0126 | 0.05 |
| BT-RM–Nash                  | 3     | 0.7375 | 0.8656 | 0.4333 | 0.0387 | 0.5458 | 0.5320 | 7.69 |
| <i>sample SD over seeds</i> |       | 0.0115 | 0.0042 | 0.0160 | 0.0046 | 0.0236 | 0.0108 | 0.07 |
| Game-utilitarian            | 3     | 0.7252 | 0.8625 | 0.4021 | 0.0389 | 0.5685 | 0.5363 | 7.66 |
| <i>sample SD over seeds</i> |       | 0.0059 | 0.0034 | 0.0130 | 0.0023 | 0.0123 | 0.0196 | 0.16 |
| Global game-maxmin          | 3     | 0.7437 | 0.8678 | 0.3917 | 0.0388 | 0.5636 | 0.5323 | 7.60 |
| <i>sample SD over seeds</i> |       | 0.0043 | 0.0024 | 0.0065 | 0.0046 | 0.0169 | 0.0203 | 0.11 |
| Scalarized DPO (uniform)    | 1     | 0.6359 | 0.8582 | 0.6281 | 0.0000 | 0.7615 | 0.7910 | 8.04 |
| PROSPER (adapt.)            | 2     | 0.7394 | 0.8654 | 0.3953 | 0.0402 | 0.5441 | 0.5215 | 7.51 |
| <i>sample SD over seeds</i> |       | 0.0052 | 0.0102 | 0.0022 | 0.0053 | 0.0035 | 0.0120 | 0.13 |
| MOPO (adapt.)               | 3     | 0.7579 | 0.8701 | 0.3156 | 0.0600 | 0.5170 | 0.5450 | 7.64 |
| <i>sample SD over seeds</i> |       | 0.0151 | 0.0016 | 0.0108 | 0.0069 | 0.0031 | 0.0330 | 0.03 |

outcomes closest to the base on the safety measures, but this alone does not establish smaller policy movement or identify an effect of its constraints. Longer DPO responses are descriptive, not a demonstrated explanation. Appendix G.1 supplies protocols, coverage and historical SafeRLHF results.

**Solver, realization, and cost.** Recomputation from stored UF tensors gives a dual residual of  $3.74 \times 10^{-14}$  and maximum inner stationarity spread  $1.88 \times 10^{-10}$ . Across three policy seeds, held-out neural nMSE is  $0.5861 \pm 0.0033$ , compared with 1 for the unchanged policy. This measures fitting to the frozen target, not generated-policy quality. The reported seed-42 projection takes 11.2 GPU-hours; it excludes other pipeline stages. Appendix F gives diagnostics, and Appendix G.3 separates feedback, generation, teacher, solver, training, evaluation, and tuning costs.

**Supporting helpfulness–safety comparison.** SafeRLHF supplies human helpfulness and harmlessness comparisons on the same response pairs. Its completed 250-update comparison uses 2000/500/1000 policy train/dev/held-out prompts and three policy seeds. Nash and L1-matched utilitarian have positive teacher-defined aggregate surplus, without a resolved Nash advantage (Table 22). These values are not independent policy outcomes. SafeRLHF IFEval means are 0.748 and 0.745, against a historical base value of 0.752 under a protocol still requiring reconciliation. Teacher truncation affects 26,790/84,000 candidate occurrences. The requested independent helpfulness–safety frontier remains pending; UF’s four-objective plot does not replace it.

## B.2 TARGET QUALITY AND NEURAL REALIZATION

We evaluate target quality and neural realization on 200 fixed development prompts (Table 5). An independent judge compares eight learner candidates with four reference candidates. Frozen targets, empirical source weights and fitted-policy weights share these candidates; fresh generations are evaluated separately. This distinguishes categorical target quality from deployment outcomes (Appendix H).

Of 60,800 rubric-order verdicts, 60,252 parse (99.10%). Requiring all 76 verdicts per prompt-objective leaves 143/131/144/167 complete prompts and 78 common to all objectives. Categorical comparisons use this shared set; fresh rows have different valid counts. Complete-case selection can affect outcomes, so missing judgments are also bounded in  $[0, 1]$  over all 200 planned prompts.

On the 78 common prompts, the exact NBPO target improves over the empirical source by  $+0.0271$   $[+0.0023, +0.0518]$  in IF and  $+0.0505$   $[+0.0273, +0.0743]$  in helpfulness. Truthfulness  $(+0.0030)$

Table 5: Independent target-to-policy evaluation. Categorical rows share 78 complete prompts; fresh rows have row-specific  $N$ . All use the same judge and reference-bank protocol.  $\min_k s_k^E$  is a finite-bank surplus, estimated only for categorical rows.

| Representation                       | $N$ | IF    | Truth | Honesty | Help  | $\min_k s_k^E$ |
|--------------------------------------|-----|-------|-------|---------|-------|----------------|
| Reference on evaluation bank         | 78  | .500  | .500  | .500    | .500  | .000           |
| Empirical source on learner pool     | 78  | .5056 | .5099 | .5105   | .4937 | −.0054         |
| NBPO exact target                    | 78  | .5327 | .5129 | .5061   | .5442 | .0054          |
| Fixed-reference exact target         | 78  | .5323 | .5130 | .5068   | .5448 | .0058          |
| Utilitarian exact target             | 78  | .5316 | .5117 | .5043   | .5437 | .0036          |
| NBPO fitted pool, seed 42            | 78  | .5494 | .5369 | .5198   | .5402 | .0222          |
| Utilitarian fitted pool, seed 42     | 78  | .5500 | .5336 | .5123   | .5399 | .0163          |
| NBPO fresh responses, seed 42        | 186 | .5455 | .5460 | .5221   | .5781 | n/a            |
| Utilitarian fresh responses, seed 42 | 188 | .5158 | .5279 | .4926   | .5575 | n/a            |
| DPO fresh responses, seed 42         | 187 | .5806 | .5548 | .5171   | .6357 | n/a            |




Figure 3: Descriptive reconstruction from Table 5. Left: exact and fitted-pool means on the same 78 prompts. Right: fresh-response means on method-specific sets (186/188/187 prompts for NBPO/utilitarian/DPO). Raw interval arrays were not supplied, so no error bars or unpaired difference estimates are reconstructed. The 0.5 line is the direct-win symmetry reference, not an estimated minimum-objective baseline.

and honesty (−0.0044) intervals include zero. Its minimum finite-bank surplus rises from −0.0054 to +0.0054. For helpfulness, the reported all-prompt missing-data bounds are [0.5722, 0.5810] for NBPO and [0.4898, 0.5001] for the source. These bound unresolved verdicts on this fixed panel; they are not population confidence intervals. The corresponding surplus bounds, [−0.16, +0.18], are uninformative.

The exact NBPO target differs from fixed-reference and utilitarian targets by at most 0.0018 in direct wins, so weak mechanism separation is already visible before neural fitting. Fitted pools improve several independent objective scores and attain minimum surplus +0.0222 (NBPO) and +0.0163 (utilitarian), despite worse target KL: about 1.457 versus source KL 1.047/1.075, TV about 0.58, and log-ratio RMS above 3.2. Thus better independent quality need not mean faithful target realization.

Fresh DPO responses have higher IF, truthfulness and helpfulness means than NBPO, while NBPO has higher honesty (.5221 versus .5171). These single-seed rows and the categorical stages use unequal prompt sets; their differences do not identify a causal projection or generation effect. Common-prompt paired comparisons are required.

A separate base draw reports objective means .4789/.4700/.4700/.4706, with each interval containing .5, and a minimum statistic .4613 [.4296, .4917]. The latter differs from the minimum of those means (.4700), so its aggregation and coverage require reconciliation. It is not a universal no-change level or a common offset to subtract from other policies. The 32 byte-identical response pairs average exactly .5; 13,312/29,893 pairs change verdict with presentation order, motivating order averaging. Finite-bank surpluses use the prespecified  $m_{\text{eval}} = 4$  and  $\beta^E = .25$  and are not population game values.

Table 6: Seed-42 projection pilot with identical pair rows and token budgets. nMSE and development pool KL measure fitting; fresh  $W_{\min}$  uses independent judgments. Common-prompt paired differences are reported in the text.

| Projection     | Steps | Train nMSE | Dev nMSE | Pool KL | Fresh $W_{\min}$ | Tokens    | GPU-h |
|----------------|-------|------------|----------|---------|------------------|-----------|-------|
| Sampled pairs  | 200   | .1665      | .9519    | 1.3112  | .5055            | 6,597,080 | 1.07  |
| All candidates | 200   | 3.5157     | .8282    | 1.4054  | .5121            | 6,597,080 | 0.93  |

The projection pilot holds targets fixed and changes within-update pair coverage (Table 6). Both seed-42 runs process the same 5,600 rows—all 28 unordered pairs per training prompt—in 200 updates, with 6,597,080 non-padding response tokens each. The sampled arm mixes prompts within an update; the all-candidate arm uses one prompt’s complete pair graph. The all-pair identity in Eq. (47) was checked in value and gradient to  $2.8 \times 10^{-14}$ . The implemented comparison changes batching; it does not benchmark the identity’s one-likelihood-per-candidate implementation.

Complete coverage improves development nMSE (.8282 versus .9519), correlation (.42 versus .23), and sign agreement (.60 versus .56), but overshoots on training data: prediction RMS 13.5 versus target RMS 5.5, with nMSE 3.5157. The sampled arm has train nMSE .1665 and prediction RMS 1.6. Repeated candidate appearances under an averaged loss do not by themselves imply a larger gradient; the effect of batch correlation and learning rate remains untested. Neither arm improves development pool KL over the source’s approximately 1.05.

On 157 shared fresh-response prompts, all-candidate minus sampled-pair differences are  $+.0167 / -.0032 / -.0008 / +.0056$  across the four objectives and  $+.0044$  for the minimum. All paired intervals include zero; the minimum interval is  $[-.0366, +.0426]$ . The pilot therefore establishes a development-fit difference, without a resolved fresh-response gain, and does not replace the 10k-prompt, 1,250-update results.

## C CONTROLLED BENCHMARK: CONSTRUCTION AND FULL RESULTS

Figure 1 in the main text summarizes this benchmark. The construction, every method, and robustness checks are retained below.

**Feasibility-preserving construction.** For every seed we form a transitive base matrix and an odd-cardinality circulant skew-symmetric perturbation. Zero row sums keep the uniform reference disagreement invariant, while bisection selects the cycle magnitude so that  $\rho^* = \max_{\pi} \min_k s_k(\pi)$  stays close to a target margin. We reject invalid probabilities rather than clip matrix entries. The resulting panel increases BT deviance from approximately  $2.7 \times 10^{-4}$  to 0.351 while retaining positive joint improvement in every instance. All policies are evaluated under the same adaptive-game values, not their method-native scales.

**Robustness across construction and temperature.** The primary construction is checked against an independently generated projected skew-symmetric family,  $C = N(R - R^T)N$  with  $N$  the symmetric projector annihilating both  $\mu$  and a witness policy, and against a temperature sweep on that family. The trend replicates: adaptive-game rows stay flat while fixed-reference and scalar-reward rows fall. The effect *size* does not: the  $\alpha = 1$  gap is  $+0.111$  on the projected family against  $+1.172$  on the circulant one, because the projected family’s BT deviance only reaches 0.174 against 0.351. We therefore claim robustness of the trend and do not attach the circulant effect size to it. Intervals in Table 9 are exact paired Student- $t$  intervals over seeds, not bootstrap intervals.

## D NATURAL SUPERVISION AND PREFERENCE MODELS

**Dataset audit.** Table 10 records the empirical role of each released dataset. ULTRA FEEDBACK supplies complete per-completion ordinal attributes, so score-induced pairwise labels are transitive by construction. SAFERLHF supplies direct helpfulness and safety comparisons but almost exclusively isolated edges; zero observed triangles means the data cannot test within-objective cycles. Its cross-

Table 7: Full controlled results over five seeds under the common adaptive-game evaluator. Residuals are method-specific; the legacy map is a non-convergence ablation.

| method                                                     | min surplus | sd     | min $s/\rho^*$ | sd exploitability | inner resid.        |
|------------------------------------------------------------|-------------|--------|----------------|-------------------|---------------------|
| $\alpha = 0.00, \rho^* = 0.0300, BT \text{ dev.} = 0.0003$ |             |        |                |                   |                     |
| NBPO direct                                                | +0.0264     | 0.0011 | 0.881          | 0.037             | $2 \times 10^{-12}$ |
| Game-KS                                                    | +0.0268     | 0.0005 | 0.893          | 0.017             | $10^{-9}$           |
| Game-util.                                                 | +0.0260     | 0.0014 | 0.866          | 0.048             | $3 \times 10^{-10}$ |
| Fixed-ref. Nash                                            | +0.0265     | 0.0011 | 0.883          | 0.036             | 0                   |
| BT-RM-Nash                                                 | +0.0265     | 0.0011 | 0.883          | 0.036             | 0                   |
| BT-RM-util.                                                | +0.0260     | 0.0014 | 0.867          | 0.047             | 0                   |
| legacy $R$ -step <sup>†</sup>                              | +0.0252     | 0.0012 | 0.840          | 0.041             | $5 \times 10^{-10}$ |
| $\alpha = 0.25, \rho^* = 0.0300, BT \text{ dev.} = 0.0202$ |             |        |                |                   |                     |
| NBPO direct                                                | +0.0269     | 0.0013 | 0.897          | 0.044             | $10^{-7}$           |
| Game-KS                                                    | +0.0272     | 0.0007 | 0.905          | 0.025             | $10^{-7}$           |
| Game-util.                                                 | +0.0265     | 0.0017 | 0.883          | 0.058             | $4 \times 10^{-8}$  |
| Fixed-ref. Nash                                            | +0.0272     | 0.0012 | 0.907          | 0.040             | 0                   |
| BT-RM-Nash                                                 | +0.0272     | 0.0012 | 0.907          | 0.040             | 0                   |
| BT-RM-util.                                                | +0.0269     | 0.0015 | 0.895          | 0.049             | 0                   |
| legacy $R$ -step <sup>†</sup>                              | +0.0258     | 0.0014 | 0.860          | 0.046             | $7 \times 10^{-1}$  |
| $\alpha = 0.50, \rho^* = 0.0300, BT \text{ dev.} = 0.0813$ |             |        |                |                   |                     |
| NBPO direct                                                | +0.0274     | 0.0013 | 0.914          | 0.044             | $2 \times 10^{-6}$  |
| Game-KS                                                    | +0.0276     | 0.0007 | 0.921          | 0.023             | $2 \times 10^{-7}$  |
| Game-util.                                                 | +0.0270     | 0.0017 | 0.900          | 0.057             | $2 \times 10^{-7}$  |
| Fixed-ref. Nash                                            | +0.0243     | 0.0018 | 0.809          | 0.059             | 0                   |
| BT-RM-Nash                                                 | +0.0243     | 0.0018 | 0.809          | 0.058             | 0                   |
| BT-RM-util.                                                | +0.0240     | 0.0019 | 0.800          | 0.064             | 0                   |
| legacy $R$ -step <sup>†</sup>                              | +0.0025     | 0.0044 | 0.082          | 0.147             | $10 \times 10^{-1}$ |
| $\alpha = 0.75, \rho^* = 0.0300, BT \text{ dev.} = 0.1882$ |             |        |                |                   |                     |
| NBPO direct                                                | +0.0278     | 0.0013 | 0.926          | 0.042             | $8 \times 10^{-6}$  |
| Game-KS                                                    | +0.0278     | 0.0005 | 0.928          | 0.018             | $8 \times 10^{-7}$  |
| Game-util.                                                 | +0.0273     | 0.0017 | 0.909          | 0.056             | $7 \times 10^{-7}$  |
| Fixed-ref. Nash                                            | +0.0138     | 0.0027 | 0.460          | 0.089             | 0                   |
| BT-RM-Nash                                                 | +0.0138     | 0.0027 | 0.460          | 0.089             | 0                   |
| BT-RM-util.                                                | +0.0137     | 0.0028 | 0.455          | 0.095             | 0                   |
| legacy $R$ -step <sup>†</sup>                              | -0.0453     | 0.0060 | -1.511         | 0.201             | $10 \times 10^{-1}$ |
| $\alpha = 1.00, \rho^* = 0.0283, BT \text{ dev.} = 0.3510$ |             |        |                |                   |                     |
| NBPO direct                                                | +0.0265     | 0.0033 | 0.931          | 0.041             | $5 \times 10^{-5}$  |
| Game-KS                                                    | +0.0264     | 0.0027 | 0.932          | 0.019             | $3 \times 10^{-6}$  |
| Game-util.                                                 | +0.0258     | 0.0037 | 0.908          | 0.063             | $2 \times 10^{-6}$  |
| Fixed-ref. Nash                                            | -0.0063     | 0.0068 | -0.241         | 0.277             | 0                   |
| BT-RM-Nash                                                 | -0.0063     | 0.0068 | -0.242         | 0.276             | 0                   |
| BT-RM-util.                                                | -0.0065     | 0.0070 | -0.248         | 0.284             | 0                   |
| legacy $R$ -step <sup>†</sup>                              | -0.0949     | 0.0064 | -3.373         | 0.406             | $10^0$              |

Table 8: Finite-pool solver audit at  $\alpha = 1$  in the original controlled family. The direct practical solve recovers the positive-surplus proximal target; the deployed legacy map does not.

| solution                | min surplus | $/\rho^*$ | inner residual     | TV to exact prox. |
|-------------------------|-------------|-----------|--------------------|-------------------|
| Exact global Nash       | +0.0158     | +0.935    | —                  | 0.007             |
| Exact proximal Nash     | +0.0158     | +0.936    | —                  | 0.000             |
| NBPO direct inner solve | +0.0139     | +0.821    | $2 \times 10^{-6}$ | 0.027             |
| Legacy $R$ -step map    | -0.1851     | -10.938   | $10^0$             | 0.643             |
| BT-RM-Nash              | -0.1001     | -5.918    | 0                  | 0.495             |

objective disagreement motivates a bargaining diagnostic; joint feasibility and separation of targets must be checked on generated responses.

#### D.1 WITHIN-OBJECTIVE CYCLIC-PREFERENCE AUDIT

Across 100 reserved prompts and 60,000 judgments, no rubric shows a repeated-cycle excess over the fitted BT noise control. No triangle satisfies the order-robust witness rule, so no confirmation judgments were collected. These results neither establish transitivity nor demonstrate cyclicity in human feedback. We test whether directly elicited preferences remain cyclic within each training

Table 9: Robustness across construction and temperature. Gaps in normalized worst surplus use paired Student- $t$  95% intervals across seeds. The temperature sweep uses the projected family.

| cycle family             | $\beta$ | seeds | NBPO—FixedRef ( $t_{95}$ ) |                    | NBPO—BT-RM ( $t_{95}$ ) |                    | min $\rho^*$ | converged |
|--------------------------|---------|-------|----------------------------|--------------------|-------------------------|--------------------|--------------|-----------|
| Circulant (primary)      | 0.25    | 5     | +1.1724                    | [+0.8715, +1.4733] | +1.1733                 | [+0.8735, +1.4731] | 0.0283       | 150/150   |
| Projected skew-symmetric | 0.25    | 5     | +0.1114                    | [+0.0714, +0.1514] | +0.1115                 | [+0.0715, +0.1514] | 0.0332       | 150/150   |
| Projected skew-symmetric | 0.10    | 3     | +0.2381                    | [+0.0880, +0.3883] | +0.2385                 | [+0.0883, +0.3887] | 0.0331       | 90/90     |
| Projected skew-symmetric | 0.50    | 3     | +0.0448                    | [+0.0165, +0.0730] | +0.0449                 | [+0.0167, +0.0730] | 0.0334       | 90/90     |

Table 10: Released-supervision audit. Zero observed human triangles means cyclicity is unidentifiable, not absent. Human conflict is cross-objective label disagreement on held-out pairs.

| dataset       | rows   | prompts | responses | obs. triangles      | human conflict | supervision        | role               |
|---------------|--------|---------|-----------|---------------------|----------------|--------------------|--------------------|
| SAFERLHF      | 73,907 | 38,639  | 143,668   | 0                   | 24.33%         | direct pairwise    | objective conflict |
| ULTRAFEEDBACK | 63,967 | 63,967  | 4/prompt  | n/a (score-induced) | n/a            | ordinal attributes | transitive control |

rubric. PROSPER reports residual cycles when LLM judges evaluate checklist items separately (Zhang et al., 2026); we adapt this comparison to our fixed objectives and add independent repeat judgments and a fitted sampling-noise control. This audit concerns new pairwise feedback on natural prompts, not cycles in UltraFeedback’s score-induced labels.

**Prompts and responses.** Reserve 100 eligible UltraFeedback prompts, stratified across Ultra-Chat, ShareGPT, and Evol-Instruct, from unused prompt groups disjoint from all five frozen training/development/final splits and the benchmark bank. Use a separate 20-prompt pilot for parser and throughput checks; exclude it from the audit. Freeze prompt IDs, rubrics, model revisions, seeds, and this protocol before viewing judgments. Generate four responses per audit prompt from the unchanged Llama-3.1-8B-Instruct reference, with temperature 0.8, top- $p = 0.9$ , and at most 2048 new tokens. Judge all six unordered pairs. Retain duplicates and truncated generations, flag both, and compare the full recorded texts without preference-based filtering. Thus the audit targets this declared candidate distribution. Extra generators or larger pools require separately labeled results.

**Objectives and repeated judgments.** Use the same written IF, truthfulness, honesty, and helpfulness rubrics as the UF-4 experiment, one rubric per query. A fifth, joint-rubric condition asks for an overall comparison and serves only as an aggregation control. Run the local Qwen3-14B judge with thinking disabled, temperature 0.8, top- $p = 0.9$ , top- $k = 20$ , and a 256-token output limit. Unlike a normalized Likert score, the primary response is an explicit choice: A, B, or tie. Require valid structured output, hide model identities, and include the complete prompt and both responses within a declared 8192-token input budget. Out-of-budget inputs are reported as missing, not silently shortened.

Collect two independent panels, A and B, each with five judgments per presentation order. Use fresh seeds and isolated contexts; correct the response orientation before analysis. The core budget is  $100 \times 6 \times 5 \times 20 = 60,000$  judgments, excluding the pilot and confirmation below. Report throughput and GPU-hours. Invalid outputs remain missing without selective retries; a complete prompt-rubric tournament requires all scheduled judgments on all six pairs. Report retained counts and missingness by cause.

**Preference graphs and metrics.** For each correctly oriented judgment, assign  $X = 1$  to a win,  $X = 0$  to a loss, and  $X = 1/2$  to a tie. For panel  $b \in \{A, B\}$ , define

$$\hat{p}_{ij}^{(b)} = \frac{1}{2} \left( \overline{X}_{ij}^{(b, i\text{-first})} + \overline{X}_{ij}^{(b, j\text{-first})} \right). \quad (40)$$

This is a tie-adjusted preference under equal presentation-order sampling, conditional on the fixed judge, rather than a human-population probability. An edge estimate requires all ten scheduled judgments; otherwise it is unresolved. Draw  $i \rightarrow j$  only when  $\hat{p}_{ij}^{(b)} > 1/2$ ; exact ties create no edge. Let  $C_A$  count prompts containing any directed cycle in panel A, including length-four cycles, and let  $R$  count those containing the *same oriented response triangle* in both panels. The latter measures repeatability, not statistical confirmation. Let  $W_A$  count prompts where every response has an observed loss; this rules out a strict winner in the observed graph. Otherwise, a graph without an

observed strict winner is unresolved. Report all three rates over the 100 planned prompts, and cyclic-triangle rates over the 400 planned triangles. Missing comparisons cannot establish an event; report unresolved counts and the corresponding worst-case upper rate, without interpreting unobserved cycles as absent. Let  $n$  be the number of tournaments complete in both panels.

**BT sampling-noise control.** A noisy transitive model can also produce observed cycles. Fit an order-aware BT model with a tie extension to panel A, with one score  $r_i$  per response, a rubric-level position coefficient  $b$ , and tie parameter  $\nu > 0$ :

$$\Pr(i, \text{tie}, j \mid o) \propto \left( e^{z/2}, 2\nu, e^{-z/2} \right), \quad z = r_i - r_j + bo, \quad o \in \{-1, +1\}. \quad (41)$$

Here  $o = +1$  presents  $i$  first; response scores sum to zero within each prompt. Fit by summed negative log likelihood with a fixed  $10^{-3}$  squared penalty on scores,  $b$ , and  $\log \nu$  to keep separated fits finite. Under equal-order averaging this null retains a transitive latent ranking. Generate 2000 synthetic A/B panel pairs under the fitted null, preserving presentation orders, repeat counts, and missingness, and rerun the graph construction. Report its mean repeated-cycle rate  $R_{\text{BT}}$ , the excess  $\Delta R = R - R_{\text{BT}}$ , and the upper-tail value  $(1 + \#\{R^{\text{sim}} \geq R\})/2001$ . These are approximate fitted-null checks, not exact composite-null tests. Apply Holm correction across the four individual rubrics; report the joint condition separately. Obtain 95% intervals with 2000 whole-prompt bootstrap replicates, retaining all judgments and refitting the null. These intervals measure prompt variability, not all uncertainty in the fitted noise model. Pairs and triangles are not independent samples.

Table 11: Cycle audit on 100 prompts.  $n$  counts tournaments complete in both panels;  $\Delta R$  is the repeated-cycle excess in percentage points over the fitted BT null, with 95% prompt-bootstrap intervals. All four Holm-adjusted checks are nonsignificant.

| Rubric        | $n$ | $C_A$ | $R$ | $R_{\text{BT}}$ | $W_A$ | $\Delta R$ [CI]     | $p_{\text{Holm}}$ |
|---------------|-----|-------|-----|-----------------|-------|---------------------|-------------------|
| IF            | 77  | 9     | 1   | 0.97            | 7     | +0.03 [−2.00, 1.00] | 1.00              |
| Truthfulness  | 73  | 2     | 0   | 0.43            | 2     | −0.43 [−2.00, 0.00] | 1.00              |
| Honesty       | 75  | 5     | 1   | 0.63            | 4     | +0.37 [−2.00, 1.00] | 1.00              |
| Helpfulness   | 96  | 4     | 0   | 0.31            | 2     | −0.31 [−2.00, 0.00] | 1.00              |
| Joint control | 90  | 1     | 0   | 0.24            | 1     | −0.24 [−1.00, 0.00] | —                 |

$C_A$ : any cycle in A;  $\bar{R}$ : the same oriented triangle in A and B;  $W_A$ : every response has an observed loss in A. Both-panel incomplete counts are 23/27/25/4/10 in row order. Unresolved cases are not cycle-free; the joint control is outside the four-rubric Holm family.

**Coverage and training relevance.** Table 12 records coverage, ties, order sensitivity, and whether the frozen GPM predicts the new comparisons better than the frozen scalar BT model. Score both on panel B using the mean per-judgment Brier loss  $(P - X)^2$  with paired prompt-bootstrap intervals; neither model is fitted or calibrated on the audit. This predictive comparison does not itself establish cycles. Label any model subsequently trained on the new judgments as a separate experiment. A planned SafeRLHF extension would repeat the audit under helpfulness and harmlessness; a claim about human preferences additionally requires fresh human comparisons on complete response triples under those same rubrics. Human annotations, prompted-judge choices, and teacher predictions must remain separate evidence sources.

Table 12: Audit coverage and prediction. Order gap measures presentation-order sensitivity; negative GPM–BT Brier differences favor GPM. All paired 95% intervals include zero; neither model was fitted or calibrated on this audit.

| Rubric        | Missing (%) | Ties (%) | Triangles (%) | Order gap | BT Brier | GPM–BT [CI]                |
|---------------|-------------|----------|---------------|-----------|----------|----------------------------|
| IF            | 0.60        | 20.8     | 3.25          | 0.535     | 0.2077   | −0.0016 [−0.0075, +0.0046] |
| Truthfulness  | 0.48        | 64.9     | 0.50          | 0.125     | 0.0999   | −0.0038 [−0.0087, +0.0012] |
| Honesty       | 0.36        | 29.3     | 1.25          | 0.305     | 0.1933   | −0.0029 [−0.0100, +0.0037] |
| Helpfulness   | 0.12        | 6.9      | 1.00          | 0.695     | 0.2341   | −0.0030 [−0.0079, +0.0014] |
| Joint control | 0.41        | 3.9      | 0.25          | 0.692     | n/a      | n/a                        |

Missing: fraction of scheduled judgments unavailable; ties: fraction of valid judgments labeled tie; triangles: observed cyclic triangles in A divided by 400; order gap: mean absolute difference between the two correctly oriented order-specific means in A, on complete pairs. Also release unresolved counts, missingness causes, and duplicate/truncation counts. Brier comparisons use identical complete panel-B pairs.

**Independent confirmation of a cycle.** Before opening B, select at most one directed cyclic triangle per individual rubric from A. Require complete A judgments and all three cycle directions to agree across both presentation orders. Among eligible cycles, maximize the minimum edge margin, breaking ties by prompt and response IDs. Fix its orientation and allocate 200 independent B judgments per selected edge, 100 per order. This adds at most  $4 \times 3 \times 190 = 2280$  judgments; only the first five B judgments per presentation order enter the aggregate tables and their null simulations. For each selected edge  $e$ , use the full  $m = 200$  bounded observations to test its order-averaged population mean against  $1/2$ :

$$p_e = \exp \left[ -2m (\max\{\hat{p}_e - 1/2, 0\})^2 \right], \quad p_{\text{cycle}} = \max_{e \text{ in triangle}} p_e. \quad (42)$$

Hoeffding’s bound gives a conservative edge test under independent judgments (Hoeffding, 1963). The cycle null states that at least one of its three edges has mean at most  $1/2$ , so the maximum edge  $p$ -value tests that union null. Apply Holm at 0.05 across the four rubric witnesses, assigning  $p = 1$  when no witness is selected or confirmation is incomplete. This establishes a cycle only for the specified judge’s order-averaged preference distribution. Require no adaptively chosen extra samples after viewing B. Report both order-specific means, and release the prompt, rubric, response texts, and individual judgments.

Table 13: Cycle-confirmation outcome. No individual rubric produced an eligible order-robust witness; no additional judgments were collected, and the protocol assigns  $p = 1$ . This does not establish transitivity.

| Rubric       | Triangle ID   | $m$ | $\hat{p}_{ab}$ | $\hat{p}_{bc}$ | $\hat{p}_{ca}$ | $p_{\text{Holm}}$ | Confirmed?    |
|--------------|---------------|-----|----------------|----------------|----------------|-------------------|---------------|
| IF           | none selected | —   | —              | —              | —              | 1.00              | Not confirmed |
| Truthfulness | none selected | —   | —              | —              | —              | 1.00              | Not confirmed |
| Honesty      | none selected | —   | —              | —              | —              | 1.00              | Not confirmed |
| Helpfulness  | none selected | —   | —              | —              | —              | 1.00              | Not confirmed |

Triangle ID resolves the prompt and response IDs ( $a, b, c$ ). Add the full prompt, rubric, and three response texts for each selected case. A missing candidate is reported as “none selected”; it is never replaced using B.

**Interpretation.** The primary witness tests select no eligible triangle. Across the four individual rubrics, panel A contains 24 cyclic triangles, but only 15/72 constituent edges retain their direction in both presentation orders. This fails the witness rule without proving that order-averaged cycles are artifacts. Secondary fitted-null checks detect no repeated-cycle excess, and the frozen GPM–BT Brier intervals all include zero. These outcomes do not establish transitivity, human-preference cyclicity, or an advantage of Nash aggregation.

## D.2 PREFERENCE-MODEL DIAGNOSTICS

**Anti-symmetric general preference model.** For objective  $k$ , a joint encoder produces  $a_{\phi,k}(x, y, z)$  and we set

$$\ell_{\phi,k}(x, y, z) = \frac{1}{2} \{a_{\phi,k}(x, y, z) - a_{\phi,k}(x, z, y)\}, \quad P_{\phi,k}(y \succ z \mid x) = \sigma(\ell_{\phi,k}(x, y, z)).$$

This guarantees exact swap complementarity and self-ties without constraining the logit to a scalar reward difference. In the SafeRLHF diagnostic, both the joint model and the scalar BT baseline use RoBERTa-base so that encoder capacity and training cost are matched; this backbone choice is an experimental control rather than part of the NBPO method. Because the observed SafeRLHF graph contains no triangles, their near-null comparison is a calibration result rather than evidence that human preferences are transitive.

## D.3 GENERAL-INSTRUCTION DATA AND PREFERENCE MODELS

The UltraFeedback panel retains UltraChat, ShareGPT, and Evol-Instruct, excludes TruthfulQA, FalseQA, and FLAN, and removes benchmark overlaps before splitting by normalized prompt identity. The realized split is 20k/2k preference-model train/dev, 10k/1k policy train/dev, and 2k final evaluation; final prompts are excluded from model selection. Filtering retained 38,240 eligible groups (Table 16), so the predeclared 5k policy-training fallback for 30k–35k eligible groups, excluding 35k,

Table 14: Held-out validation of checkpointed GPM and BT ensembles, seeds 41/42/43. Calibration uses validation data. Sparse comparison graphs cannot establish a cyclicity advantage.

| model              | objective   | per-seed (mean of 3) |          |        | $T$   | calibrated 3-seed ensemble |          |        |        |         |
|--------------------|-------------|----------------------|----------|--------|-------|----------------------------|----------|--------|--------|---------|
|                    |             | NLL                  | bal. acc | AUC    |       | NLL                        | bal. acc | AUC    | ECE    | ens. sd |
| Anti-symmetric GPM | helpfulness | 0.5543               | 0.7218   | 0.7999 | 1.282 | 0.5512                     | 0.7257   | 0.8034 | 0.1228 | 0.0293  |
| Anti-symmetric GPM | harmless.   | 0.5180               | 0.7556   | 0.8464 | 1.297 | 0.5131                     | 0.7573   | 0.8512 | 0.1782 | 0.0357  |
| Scalar BT          | helpfulness | 0.5560               | 0.7214   | 0.7974 | 1.270 | 0.5527                     | 0.7250   | 0.8009 | 0.1220 | 0.0300  |
| Scalar BT          | harmless.   | 0.5212               | 0.7507   | 0.8431 | 1.265 | 0.5163                     | 0.7544   | 0.8484 | 0.1779 | 0.0357  |

Table 15: Generated-pool pilot on 200 development prompts. Conflict and cycles are model predictions, not human-label measurements; cycles use a common comparator tournament.

| geom. | dup. rate | entropy | ens. sd | GPM/BT sign | model conflict | pred. cycles | split-half $r$ | $\rho$ | sign agr. | dual evals | solve (s) |
|-------|-----------|---------|---------|-------------|----------------|--------------|----------------|--------|-----------|------------|-----------|
| 4+4   | 0.0073    | 0.6117  | 0.0429  | 0.9102      | 7.67%          | 0/800        | 0.9172         | 0.9132 | 0.8667    | 42         | 24.8      |
| 8+8   | 0.0167    | 0.6120  | 0.0414  | 0.9061      | 7.80%          | 10/11200     | 0.9242         | 0.9180 | 0.8925    | 30         | 18.4      |

was not used; fewer than 30k would require revising the budget. Honesty was the binding attribute: 1,494 groups lacked two valid ratings, versus at most one group for each other attribute. Split hashes are recorded in the experiment artifacts.

Released GPT-4 ratings train four attribute heads: valid comparisons receive 0/1 labels, ties receive 0.5, and missing labels are masked. Four completions yield six unordered pairs per prompt. These rating-derived labels are transitive by construction and cannot establish preference cycles. Policy-training prompts are disjoint and use newly sampled base-policy responses.

The GPM and scalar BT model share the 149M ModernBERT-base encoder with native 8192-token support (Warner et al., 2024), supervision, prompt splits, and one preference-model seed. The GPM uses the antisymmetric construction above; BT uses differences of four scalar scores. Training uses AdamW, learning rate  $2 \times 10^{-5}$ , batch 32 pairs, weight decay 0.01, warmup 0.06, length bucketing, and up to two epochs, selected on preference-model development data. Policy-seed repeats therefore do not measure teacher uncertainty. Unsupported inputs are audited; an independent judge evaluates full generated responses.

Policy training uses full 8B fine-tuning, batch 32 pairs, the existing peak learning-rate scale, and 1250 updates for the realized 10k training prompts; the unused 5k fallback specifies 625 updates. This gives four sampled pair exposures per prompt on average. Each run has its own cosine schedule and one outer stage. Independent learner and comparator pools contain eight occurrences each; reference-reference scores and the shared tensor are cached across methods. Two encoder orders and four joint heads require  $2 \times 10^4(64 + 28) = 1.84$  million encoder passes for 10k prompts.

The 20.00% generation-cap rate counts responses reaching the 1024-token horizon; the 0/500,000 scoring-oversize count concerns serialized encoder inputs. These have different denominators. Reference-tensor antisymmetry and zero diagonal are imposed by construction and do not validate the teacher.

Table 17 reports neural fitting diagnostics. For comparison, SafeRLHF achieved nMSE  $.6845 \pm .0048$  and sign accuracy  $.7009 \pm .0016$  over three seeds, with target RMS 1.125; normalized errors should be interpreted alongside these different target scales. Nash and utilitarian targets have total variation .0195, and the two arms attain similar fitting accuracy. Neither observation establishes a policy-quality difference.

The pilot showed substantial first-position skew (Table 19); the audit evaluates both presentation orders and reports their average. Capability measurements in Table 18 are separate from the four alignment objectives.

**Aggregation across objectives and across prompts.** NBPO applies  $\sum_k \log(V_k(\pi) - d_k)$  to prompt-averaged objective values. In contrast, PROSPER’s regularized criterion has prompt-dependent

Table 16: Reported UF-4 data, preference-model, pool, and solver measurements. IF, Truth, Honesty, and Help denote instruction-following, truthfulness, honesty, and helpfulness.

| Stage       | Quantity                                                | Value                        |
|-------------|---------------------------------------------------------|------------------------------|
| Data        | Raw rows: UltraChat / ShareGPT / Evol-Instruct          | 9,929/19,949/10,000          |
|             | Exact normalized prompts                                | 39,846                       |
|             | Groups after verified near-duplicate merge              | 39,743                       |
|             | Benchmark-overlap groups removed (all AlpacaEval)       | 9                            |
|             | Groups overlapping SafeRLHF                             | 0                            |
|             | Eligible groups: two rated completions per objective    | 38,240                       |
|             | Realized regime                                         | target sizes, no fallback    |
|             | Eligible groups left unassigned                         | 3,240                        |
|             | Groups with fewer than two honesty ratings              | 1,494                        |
|             | Preference-model pairs, train / dev                     | 120000/12000                 |
| Supervision | Tie fraction, IF / Truth / Honesty / Help               | 0.253/0.481/0.308/0.262      |
|             | Masked train pairs (honesty only)                       | 3,668                        |
|             | GPM selected step / mean dev NLL / worst NLL            | 3750/0.5601/0.6305           |
| Teacher     | GPM dev accuracy, non-ties: IF / Truth / Honesty / Help | 0.7958/0.7469/0.7677/0.8346  |
|             | BT selected step / mean dev NLL / worst NLL             | 3500/0.5589/0.6283           |
|             | BT dev accuracy, non-ties: IF / Truth / Honesty / Help  | 0.7995/0.7527/0.7599/0.8403  |
|             | Oversize / total serialized training pair sequences     | 216/240,000                  |
|             | Response truncations / maximum serialized tokens        | 0/4,323                      |
| Pool        | Sampled occurrences (learner 8 + comparator 8)          | 160,000                      |
|             | Responses stopped by the 1024-token cap                 | 20.00%                       |
|             | Duplicate occurrences retained with multiplicity        | 725                          |
| Scoring     | GPM encoder examples, realized / predicted              | 1,840,000/1,840,000          |
|             | Aggregate throughput on four H200s                      | 577.6 sequences/s            |
|             | Oversize / total serialized scoring sequences           | 0/500,000                    |
| Solve       | Reference-tensor skew / diagonal residual               | 0.0/0.0                      |
|             | Nash $\lambda$ , IF / Truth / Honesty / Help            | 7.4627/11.0426/8.1138/7.3043 |
|             | Nash $\lambda$ $L_1$                                    | 33.9235                      |
|             | Nash train surplus, IF / Truth / Honesty / Help         | 0.1340/0.0906/0.1232/0.1369  |
|             | Independently computed Nash dual residual, train        | $3.74 \times 10^{-14}$       |
|             | Maximum inner stationarity spread over prompts          | $1.88 \times 10^{-10}$       |

Table 17: UF-4 neural projection diagnostics on held-out targets. These measure target realization, not the quality of generated responses; three-seed summaries use sample standard deviation.

| Run                      | Quantity                                                             | Value                        |
|--------------------------|----------------------------------------------------------------------|------------------------------|
| NBPO-MSE, seed 42        | Optimizer updates, completed / planned                               | 1,250/1,250                  |
|                          | Held-out dev pairs per evaluation                                    | 28,000                       |
|                          | Dev nMSE $\downarrow$ , first $\rightarrow$ last of five evaluations | 0.6083 $\rightarrow$ 0.5823  |
|                          | Dev sign accuracy $\uparrow$ , first $\rightarrow$ last              | 0.6605 $\rightarrow$ 0.6694  |
|                          | Final dev Pearson / Spearman                                         | 0.6515/0.5479                |
|                          | Dev target RMS / fitted log-ratio RMS                                | 5.4711/3.8409                |
|                          | Wall time, four H200s / GPU-hours                                    | 2.79 h / 11.2 GPU-h          |
|                          | Optimizer updates, completed / planned                               | 1,250/1,250                  |
|                          | Dev nMSE $\downarrow$ , first $\rightarrow$ last of five evaluations | 0.6061 $\rightarrow$ 0.5862  |
|                          | Dev sign accuracy $\uparrow$ , first $\rightarrow$ last              | 0.6603 $\rightarrow$ 0.6702  |
| Utilitarian, seed 42     | Final dev Pearson / Spearman                                         | 0.6488/0.5480                |
|                          | Dev target second moment, Nash / utilitarian                         | 29.9332/31.2487              |
|                          | Final dev nMSE $\downarrow$ , mean $\pm$ SD                          | 0.5861 $\pm$ 0.0033          |
|                          | Final dev sign accuracy $\uparrow$ , mean $\pm$ SD                   | 0.6698 $\pm$ 0.0017          |
|                          | Final dev Pearson, mean $\pm$ SD                                     | 0.6490 $\pm$ 0.0023          |
| Utilitarian, three seeds | Final nMSE, seeds 42 / 43 / 44                                       | 0.5823/0.5881/0.5878         |
|                          | Final dev nMSE $\downarrow$ , mean $\pm$ SD                          | 0.5800 $\pm$ 0.0054          |
|                          | Final dev sign accuracy $\uparrow$ , mean $\pm$ SD                   | 0.6732 $\pm$ 0.0059          |
|                          | Final nMSE, seeds 42 / 43 / 44                                       | 0.5862/0.5778/0.5761         |
|                          | Train $\lambda$ (IF / Truth / Hon / Help)                            | 7.6311/11.1646/8.2602/7.4579 |
| Fixed-reference Nash     | Train $\lambda$ $L_1$ , fixed-reference / adaptive                   | 34.5138/33.9235              |
|                          | Train minimum surplus, fixed-reference / adaptive                    | 0.08957/0.09056              |
|                          | Dual box active                                                      | no                           |

weights (Zhang et al., 2026):

$$\min_{w: \mathcal{X} \rightarrow \Delta_K, \nu} \mathbb{E}_x \left[ \sum_k w_k(x) J_{k,x}(\pi, \nu) + \beta \text{KL}(\nu(\cdot | x) \| \pi_{\text{ref}}(\cdot | x)) \right],$$

where  $J_{k,x}$  is the conditional pairwise payoff. Constant shifts from centered payoffs do not change its optimizer. Prompt-dependent weights cannot be replaced by a shared weight without changing this

Table 18: UF-4 capability measurements under the shared evaluation recipe. These point estimates do not establish equivalence; trained-policy IFEval and safety measurements remain pending.

| Checkpoint(s)             | Benchmark / metric                        | Value                |
|---------------------------|-------------------------------------------|----------------------|
| Base                      | ARC-Challenge 25-shot, acc / acc_norm     | 0.5674/0.6049        |
|                           | HellaSwag 10-shot, acc / acc_norm         | 0.5927/0.8008        |
|                           | IFEval strict-prompt / strict-instruction | 0.7542/0.8249        |
|                           | GSM8K exact match                         | 0.8666               |
|                           | HarmBench harmful-completion rate         | 0.2812               |
| Base / NBPO / utilitarian | XSTest 450                                | unadjudicated        |
|                           | ARC-Challenge 25-shot, acc_norm           | 0.6049/0.6058/0.6032 |
|                           | HellaSwag 10-shot, acc_norm               | 0.8008/0.7984/0.7978 |
| Base / NBPO               | MMLU 5-shot, acc (utilitarian pending)    | 0.6823/0.6816        |

Table 19: UF-4 audit execution diagnostics. The 20 pilot prompts are excluded from audit estimates; no candidate cycle met the confirmation eligibility rule.

| Stage        | Quantity                                       | Value             |
|--------------|------------------------------------------------|-------------------|
| Pilot        | Judgments collected / planned                  | 12,000/12,000     |
|              | Explicit-verdict parse rate                    | 0.9959            |
|              | Parse failures: no verdict / multiple verdicts | 40/9              |
|              | Throughput on one H200 / mean output tokens    | 33.3 per s / 99.5 |
|              | First-shown verdicts: A / B / tie              | 6,943/2,345/2,663 |
|              | Judge inputs truncated at 8192 tokens          | 0                 |
| Confirmation | Individual-rubric cyclic triangles in panel A  | 24                |
|              | Cycle edges holding in both orders             | 15/72 = 0.208     |
|              | Eligible witness triangles                     | 0 of 4 rubrics    |
|              | Judgments spent / confirmation budget          | 0/2,280           |

criterion. Prompt-direct optimization decomposes only the inner computation; NBPO’s  $\lambda$  remains shared across prompts.

To illustrate aggregation order alone, take two equally likely prompts. Surplus vectors  $(.2, 0)$  and  $(0, .2)$  average to  $(.1, .1)$ , whereas  $(.08, .08)$  on both prompts averages to  $(.08, .08)$ . Global Nash prefers the first policy, while the average prompt-wise minimum prefers the second. This schematic example omits opponent regularization and is not a numerical PROSPER result.

The reported PROSPER adaptation uses prompt-wise absolute-maxmin aggregation at  $\beta = 0.25$  and matches coefficient mass to the Nash dual (33.92). The source reports certified solutions on 11,000 train/dev prompts and a maximum log-ratio identity residual of  $9.9 \times 10^{-11}$ . Training objective-weight means are (6.97, 12.03, 7.56, 7.36) with SDs (1.93, 4.09, 1.63, 3.14); the mean effective number of active objectives is 3.50. These measurements distinguish prompt-dependent weights from a shared vector, but do not by themselves establish a faithful reproduction of every published training step.

The selected target has a negative minimum per-prompt surplus on 642/10,000 training prompts (6.42%) and 49/1,000 development prompts (4.90%). These are selected-policy violation rates, not infeasibility certificates. A different policy may jointly improve the objectives; certifying otherwise requires a nonpositive upper bound from the separate maximum-margin problem in Appendix A.5. Aggregate and per-prompt improvement remain distinct. Generated-policy results are reported in Table 1.

## E TRAINING PROTOCOL AND BASELINES

The completed Nash/utilitarian comparison uses one outer stage and policy seeds 42/43/44. Earlier WBC and long-schedule diagnostics each use one policy seed. Comparisons within the repaired panel share the stated prompts, teacher, pools, and realization settings. Equal update counts do not imply equal FLOPs. Table 20 records the original recipe; resolved configurations and external-baseline results still require source verification.

Table 20: Reported study dimensions and announced repair settings. Optimizer, precision, and sampling settings must be reconciled with the execution manifest. This table concerns the SafeRLHF panel; UltraFeedback settings appear in Appendix D.3.

| Setting                                  | Value                                                  |
|------------------------------------------|--------------------------------------------------------|
| Base / adaptation                        | Llama-3.1-8B-Instruct / full fine-tuning               |
| Train / dev / held-out prompts           | 2,000/500/1,000                                        |
| Policy seeds / outer stages              | 42/43/44 (Nash/utilitarian); 1 stage                   |
| Learner / comparator occurrences         | 8/8; independent streams                               |
| Sampling                                 | temperature 1, top- $p = 1$ ; preserve terminal tokens |
| Preference models                        | frozen three-seed objective ensembles                  |
| Inner solve / neural target              | direct finite-pool / canonical log-ratio               |
| Optimizer / learning rate                | AdamW / $5 \times 10^{-7}$                             |
| Adam moments / epsilon                   | (0.9, 0.95) / $10^{-8}$                                |
| Schedule / declared updates              | cosine / 1750 or 250                                   |
| Warmup ratio / weight decay              | 0.10/0                                                 |
| Global batch / pair rows                 | 32 pairs / 56,000                                      |
| Prompt / total sequence limit            | 1024/2048 tokens                                       |
| Precision / gradient clip                | bf16 forward, fp32 reduction / 1                       |
| Opponent temperatures / proximal $\eta$  | — (resolved config required)                           |
| Reference cache / optimizer-state checks | — (execution report required)                          |
| Model / split / config revisions         | — (immutable manifest required)                        |

Table 21: Matched mechanism controls. External PROSPER and MOPO comparisons require their actual algorithms and separate tuning contracts; game-maxmin or WBC alone is not their reproduction.

| method               | within-objective representation | aggregation      | changed component relative to NBPO            |
|----------------------|---------------------------------|------------------|-----------------------------------------------|
| NBPO                 | adaptive regularized game       | Nash             | —                                             |
| Fixed-reference Nash | fixed reference margin          | Nash             | removes adaptive opponent                     |
| BT-RM–Nash           | scalar BT reward                | Nash             | imposes scalar reward representation          |
| Game-utilitarian     | adaptive regularized game       | equal-weight sum | replaces Nash aggregation                     |
| Global game-maxmin   | adaptive regularized game       | global minimum   | replaces Nash aggregation                     |
| BT-RM–utilitarian    | scalar BT reward                | equal-weight sum | scalar representation and aggregation control |

## F SOLVER AND NEURAL-REALIZATION DIAGNOSTICS

**Completed SafeRLHF teacher-fit comparison.** The completed study uses 2000 policy-training prompts, 500 development prompts, 1000 held-out prompts, and the fixed 250-update regression recipe. Both Nash and L1-matched utilitarian attain positive teacher-defined aggregate surplus across three seeds; their paired objective intervals include zero. The schedule and pool diagnostics below concern target realization and do not replace independent policy evaluation.

The direct solve concerns a finite pool at a declared tolerance. It does not make the preference model, sampling approximation, or neural fit exact. Solver consistency, held-out regression, and generated-policy evaluation are reported separately.

**Alternative realization.** For comparison with (26), weighted behavior cloning uses

$$\mathcal{L}_{\text{WBC}}^{(t)}(\theta) = \mathbb{E}_x \sum_{i=1}^{n_x} p_{t,x}^*(i) [-\log \pi_\theta(y_{x,i} | x)]. \quad (43)$$

Weights are detached and full-vocabulary response log-probabilities are summed over tokens. With IID reference-policy occurrences, this weighted likelihood approximates population forward-KL projection; it remains a finite-sample realization rather than an exact population update (Agnihotri et al., 2026). Averaging  $\frac{n_x}{2} [p_i^*(-\log \pi_\theta(y_i | x)) + p_j^*(-\log \pi_\theta(y_j | x))]$  over all unordered pairs equals the prompt-wise objective. Uniform teacher weights need not produce zero gradient on a finite sampled dataset.

**Metrics and selection history.** Write  $t = \hat{h}_t^*$  and  $h = h_t(\pi_\theta)$ . We report  $\text{nMSE} = \mathbb{E}[(h - t)^2] / \mathbb{E}[t^2]$ , sign agreement, Pearson and Spearman correlations. Expectations use the stated split and

Table 22: Completed SafeRLHF comparison, three policy seeds per method. Entries reproduce reported means and spreads; the spread convention and minimum/averaging order require artifact verification. Surplus uses the training teachers and does not certify population individual rationality.

| Family                   | Seeds | nMSE↓               | Sign↑               | $\min_k s_k \uparrow$ |
|--------------------------|-------|---------------------|---------------------|-----------------------|
| NBPO-Nash                | 3     | $0.6845 \pm 0.0048$ | $0.7009 \pm 0.0016$ | $0.1215 \pm 0.0028$   |
| Utilitarian (L1-matched) | 3     | $0.6737 \pm 0.0242$ | $0.7040 \pm 0.0058$ | $0.1220 \pm 0.0029$   |

Table 23: Target realization on 500 development prompts, one seed. Each horizon has its own cosine schedule; 250-update runs are development-informed follow-ups. Pool KL is  $D_{\text{KL}}(p^* \parallel \hat{p}_\theta)$  on sampled support; the reference value is 0.7879. Full diagnostics are in Appendix F.

| Projection | Horizon | nMSE   | Sign  | Pearson | Pool KL↓ | GPU-h |
|------------|---------|--------|-------|---------|----------|-------|
| MSE        | 1750    | 1.418  | 0.693 | 0.415   | 1.2131   | 7.48  |
| WBC        | 1750    | 13.070 | 0.486 | -0.152  | 6.0634   | 6.78  |
| MSE†       | 250     | 0.690  | 0.702 | 0.588   | 0.6204   | 2.20  |
| WBC†       | 250     | 0.780  | 0.687 | 0.492   | 0.6875   | 2.08  |

pair measure. The zero predictor has nMSE 1. The reported oracle rescaling diagnostic is

$$a^* = \frac{\mathbb{E}[ht]}{\mathbb{E}[h^2]}, \quad \text{nMSE}^* = \min_{a \in \mathbb{R}} \frac{\mathbb{E}[(ah - t)^2]}{\mathbb{E}[t^2]}.$$

It fits a scalar on the evaluation observations and can reverse sign. It is neither centered explained variance nor a deployable calibration result. A direction-preserving version restricts  $a \geq 0$ ; signed coefficients and this additional metric require the raw moments. If  $\mathbb{E}[h^2] = 0$ , set  $a^* = 0$  and the residual to 1. Mean response log-ratio is measured under the sampled occurrence distribution; it is not total pool probability or a full-policy KL certificate.

The 1750-update comparison was specified first. The 250-update schedules were chosen after inspecting development trajectories and declared before their runs. Each uses its own warmup and cosine horizon. A checkpoint at step 250 of the long run is therefore a different condition. Diagnostic thresholds were retained, but their use in development precludes describing the entire  $2 \times 2$  as an untouched preregistered comparison.

**Finite-pool distribution fidelity.** For a candidate occurrence  $i$ , write  $\Delta_i = \log \pi_\theta(y_i \mid x) - \log \pi_t(y_i \mid x)$  and reconstruct

$$\hat{p}_\theta(i \mid x) = \frac{p_{t,x}(i)e^{\Delta_i}}{\sum_j p_{t,x}(j)e^{\Delta_j}}.$$

This self-normalized importance distribution is evaluated against the solved target by forward KL on the fixed sampled support. It does not measure the full LLM divergence or probability outside that support. MSE-250 and WBC-250 reduce KL by 0.1675 and 0.1004, respectively; the reported intervals in Table 25 concern Pearson correlation, not KL improvement. Paired prompt intervals for KL remain to be reported. With uniform occurrence weights and  $\Delta = 0$ , target-argmax mass is algebraically  $1/8$ ; reference top-1 agreement depends on tie-breaking.

**Implementation history.** The repaired study changed target construction, reference handling, tokenization, sampling, and training data jointly. Earlier failures cannot identify the effect of any single correction. Detailed debugging sweeps and the retired prompted-judge audit are retained in the code archive rather than used as performance evidence.

## F.1 AGGREGATION-TARGET SENSITIVITY

On SafeRLHF, Nash weights (6.325, 6.186) give mean target TV 0.0008 relative to equal weights with the same  $\ell_1$  mass. Pairwise-target difference RMS is 0.0034, about 0.30% of target RMS 1.125. On UF-4, mean TV is 0.0195 and relative target RMS is 3.44% (Table 27). Utilitarian attains higher unregularized  $F$  but incurs greater proximal KL, leaving Nash with higher  $J = F - D/\eta$ . Matching

Table 24: All recorded development evaluations of the repaired runs. The first column identifies the declared schedule; “Updates” is the checkpoint step. Oracle nMSE\* permits signed rescaling. No trajectory point is substituted for its run’s declared final checkpoint.

| Projection                 | Updates | nMSE   | nMSE* | Sign acc. | Pearson | Spearman | $\log(\pi_\theta/\pi_t)$ |
|----------------------------|---------|--------|-------|-----------|---------|----------|--------------------------|
| NBPO-MSE, 1750             | 250     | 0.839  | 0.734 | 0.692     | 0.516   | 0.499    | -1.13                    |
|                            | 500     | 1.226  | 0.796 | 0.688     | 0.452   | 0.465    | -1.86                    |
|                            | 750     | 1.482  | 0.824 | 0.678     | 0.419   | 0.452    | -2.42                    |
|                            | 1000    | 1.568  | 0.839 | 0.694     | 0.401   | 0.449    | -2.38                    |
|                            | 1250    | 1.517  | 0.836 | 0.688     | 0.404   | 0.450    | -2.27                    |
|                            | 1500    | 1.432  | 0.831 | 0.691     | 0.412   | 0.455    | -2.12                    |
|                            | 1750    | 1.418  | 0.828 | 0.693     | 0.415   | 0.460    | -2.07                    |
| NBPO-WBC, 1750             | 250     | 0.920  | 0.824 | 0.645     | 0.419   | 0.415    | -0.75                    |
|                            | 500     | 1.932  | 0.938 | 0.623     | 0.250   | 0.297    | -2.04                    |
|                            | 750     | 3.776  | 0.995 | 0.569     | 0.074   | 0.124    | -3.27                    |
|                            | 1000    | 6.121  | 0.999 | 0.535     | -0.028  | 0.018    | -4.39                    |
|                            | 1250    | 11.611 | 0.979 | 0.496     | -0.143  | -0.096   | -6.15                    |
|                            | 1500    | 12.672 | 0.979 | 0.489     | -0.144  | -0.097   | -6.63                    |
|                            | 1750    | 13.070 | 0.977 | 0.486     | -0.152  | -0.105   | -6.68                    |
| NBPO-MSE, 250 <sup>†</sup> | 50      | 0.737  | 0.702 | 0.675     | 0.546   | 0.507    | -0.77                    |
|                            | 100     | 0.751  | 0.674 | 0.691     | 0.571   | 0.537    | -1.11                    |
|                            | 150     | 0.710  | 0.659 | 0.695     | 0.584   | 0.548    | -1.03                    |
|                            | 200     | 0.701  | 0.668 | 0.693     | 0.576   | 0.539    | -0.87                    |
|                            | 250     | 0.690  | 0.655 | 0.702     | 0.588   | 0.552    | -0.90                    |
| NBPO-WBC, 250 <sup>†</sup> | 50      | 0.866  | 0.839 | 0.628     | 0.401   | 0.388    | -0.39                    |
|                            | 100     | 0.807  | 0.777 | 0.655     | 0.473   | 0.450    | -0.55                    |
|                            | 150     | 0.755  | 0.730 | 0.678     | 0.520   | 0.499    | -0.64                    |
|                            | 200     | 0.770  | 0.743 | 0.682     | 0.507   | 0.497    | -0.57                    |
|                            | 250     | 0.780  | 0.758 | 0.687     | 0.492   | 0.488    | -0.51                    |

Table 25: Independent reconstruction on the development pool. Correlation intervals are prompt-clustered. KL improvement is reference minus policy KL. Target-argmax mass is  $\hat{p}_\theta(i^*)$ ,  $i^* \in \arg \max_i p^*(i)$ ; tie rules must be fixed. The WBC-1750 improvement is reported at source precision and differs by  $10^{-4}$  from subtracting rounded KL entries.

| Policy                 | Pearson [clustered CI]  | $D_{\text{KL}} \downarrow$ | KL improvement <sup>†</sup> | Top-1 | Argmax mass |
|------------------------|-------------------------|----------------------------|-----------------------------|-------|-------------|
| Base                   | —                       | 0.7879                     | +0.0000                     | 0.108 | 0.1250      |
| MSE (1750)             | +0.417 [0.370, 0.463]   | 1.2131                     | -0.4252                     | 0.420 | 0.3135      |
| WBC (1750)             | -0.150 [-0.205, -0.095] | 6.0634                     | -5.2754                     | 0.124 | 0.1113      |
| MSE (250) <sup>†</sup> | +0.589 [0.554, 0.622]   | 0.6204                     | +0.1675                     | 0.406 | 0.2906      |
| WBC (250) <sup>†</sup> | +0.496 [0.459, 0.531]   | 0.6875                     | +0.1004                     | 0.362 | 0.2283      |

coefficient mass therefore does not match realized policy movement; target distance alone does not determine statistical power.

## G EVALUATION PROTOCOL

**Common-prompt denominators for the UF-4 final evaluation.** The planned final split is 2,000 prompts. A prompt enters the reported table only when every evaluated method has both presentation orders parsed on all four attributes, which leaves 1,236. Each method’s own valid count is larger, because the losses are not on the same prompts; across the evaluated arms those counts span 1,884–1,941. Reporting those per-method counts side by side would compare arms on slightly different prompts, so the common set is used throughout and the difference is recorded here. Invalid judgments are left missing; none is recoded as a tie or a loss and none is selectively retried.

### G.1 COMPLETE CAPABILITY RESULTS

This table retains the full reported capability panel. The UF and historical SafeRLHF base rows use different recorded protocols and are not pooled. The UF rows are policy seed 42. ARC-C, HellaSwag and MMLU are all complete on all three policy seeds of both families. Seed means with sample standard deviations are  $0.6064 \pm 0.0010$ ,  $0.7994 \pm 0.0011$  and  $0.6820 \pm 0.0006$  for NBPO, and  $0.6044 \pm 0.0013$ ,  $0.7991 \pm 0.0012$  and  $0.6822 \pm 0.0005$  for game-utilitarian, against base values

Table 26: Measured direct-solver scaling by prompt count, pool size, and CPU workers. Times cover finite-pool solving and artifact generation, not LLM rollout or training.

| prompts | pool  | workers | solve (s) | dual evals | s/eval | proj. KKT           | inner resid.         | peak RSS (MB) | artifact (s) |
|---------|-------|---------|-----------|------------|--------|---------------------|----------------------|---------------|--------------|
| 200     | 4+4   | 32      | 1.5       | 31         | 0.047  | $3 \times 10^{-14}$ | $3 \times 10^{-11}$  | 467           | 0.01         |
| 200     | 8+8   | 32      | 1.7       | 35         | 0.049  | $4 \times 10^{-15}$ | $5 \times 10^{-12}$  | 498           | 0.02         |
| 200     | 16+16 | 32      | 1.8       | 32         | 0.055  | $2 \times 10^{-14}$ | $10^{-11}$           | 499           | 0.02         |
| 1000    | 4+4   | 96      | 4.4       | 33         | 0.132  | $7 \times 10^{-15}$ | $2 \times 10^{-11}$  | 499           | 0.03         |
| 1000    | 8+8   | 96      | 3.2       | 26         | 0.122  | $4 \times 10^{-15}$ | $6 \times 10^{-12}$  | 504           | 0.04         |
| 1000    | 16+16 | 96      | 4.4       | 37         | 0.119  | $3 \times 10^{-13}$ | $10 \times 10^{-11}$ | 534           | 0.07         |
| 7000    | 4+4   | 96      | 10.6      | 28         | 0.379  | $7 \times 10^{-10}$ | $4 \times 10^{-11}$  | 534           | 0.11         |
| 7000    | 8+8   | 96      | 10.9      | 27         | 0.404  | $2 \times 10^{-15}$ | $3 \times 10^{-11}$  | 544           | 0.22         |
| 7000    | 16+16 | 96      | 15.3      | 30         | 0.508  | $10^{-14}$          | $7 \times 10^{-11}$  | 633           | 0.39         |

Table 27: Nash versus mass-matched utilitarian targets by panel. TV measures distribution distance; relative target RMS measures log-ratio differences. UF-4 quantities were reported as recomputed from stored arrays independently of the solver.

| Diagnostic                                       | Value                  |
|--------------------------------------------------|------------------------|
| SafeRLHF-2k: mean target TV                      | 0.0008                 |
| SafeRLHF-2k: maximum target TV                   | 0.0075                 |
| SafeRLHF-2k: pairwise target difference RMS      | 0.0034                 |
| SafeRLHF-2k: target RMS on the aggregation panel | 1.125                  |
| SafeRLHF-2k: Nash $\lambda_{\text{help}}$        | 6.325                  |
| SafeRLHF-2k: Nash $\lambda_{\text{safe}}$        | 6.186                  |
| UF-4: mean target TV                             | 0.0195                 |
| UF-4: maximum target TV                          | 0.1258                 |
| UF-4: pairwise target difference RMS             | 0.1857                 |
| UF-4: target RMS on the aggregation panel        | 5.4024                 |
| UF-4: relative pairwise target RMS               | 0.0344                 |
| UF-4: Nash $\lambda L_1$                         | 33.9235                |
| UF-4: largest / smallest Nash $\lambda$          | 11.0426/7.3043         |
| UF-4: Nash $F = \sum_k \log s_k$                 | -8.493723              |
| UF-4: utilitarian $F$                            | -8.469298              |
| UF-4: Nash realized $D(\pi    \pi_t)$            | 1.080072               |
| UF-4: utilitarian realized $D(\pi    \pi_t)$     | 1.106423               |
| UF-4: Nash $J = F - D/\eta$                      | -9.573795              |
| UF-4: utilitarian $J$                            | -9.575722              |
| UF-4: $J$ gap, utilitarian minus Nash            | -0.001927              |
| UF-4: independent Nash dual residual on train    | $3.74 \times 10^{-14}$ |

of 0.6049, 0.8008 and 0.6823. Every seed-mean change from the base is within 0.0017 (largest: game-utilitarian HellaSwag, -0.0017), no individual cell differs by more than 0.0030, and the seed spreads are 0.0005–0.0013. These are point estimates on fixed benchmarks, not equivalence tests. Table 4 in the main text isolates the core outcomes to complete on the UF checkpoints.

**Capability panel protocols.** IFEval uses strict-prompt accuracy, GSM8K exact match, HarmBench harmful-completion rate and XSTest benign refusal on 250 safe prompts; handling of its 200 unsafe prompts is reported separately. Table 4 gives the current seed counts and sample SDs. The source reports that unseeded language detection can change IFEval rescoring by up to 2/541; the base retains its frozen first scoring.

AE2 and AH2 use the AlpacaEval-2 and Arena-Hard-v2 prompt sets with a pinned local reward model (Skywork-Reward-V2-Qwen3-8B); the main AH2 column covers 500 hard prompts, separately from 250 creative prompts. MT uses 80 two-turn questions, the declared category-dependent grading templates and a 1–10 scale, with local Qwen3-14B replacing the original annotator. These are local proxy results, not official leaderboard scores. Unparsed ratings remain missing. The source’s repeated MT grading of identical responses changed the thirteen shared arm means by -0.100 to +0.056 as batch composition changed; retain this reproducibility limitation when interpreting small method differences. No equivalence claim follows from a small spread or an interval containing zero.

Table 28: Paired Nash-minus-utilitarian teacher-surplus differences on 1000 SafeRLHF prompts, conditional on each policy seed. Intervals retain the source-reported construction and all include zero; they do not establish equivalence.

| Seed | $\Delta s_{\text{help}}$ | Reported CI          | $\Delta s_{\text{safe}}$ | Reported CI          |
|------|--------------------------|----------------------|--------------------------|----------------------|
| 42   | -0.0003                  | $[-0.0067, +0.0058]$ | -0.0010                  | $[-0.0074, +0.0050]$ |
| 43   | +0.0019                  | $[-0.0038, +0.0076]$ | +0.0004                  | $[-0.0052, +0.0061]$ |
| 44   | -0.0033                  | $[-0.0086, +0.0023]$ | -0.0037                  | $[-0.0095, +0.0017]$ |

Table 29: Capability scores under each panel’s evaluation protocol. ARC-C/HellaSwag use normalized accuracy with 25/10 shots; MMLU uses 5-shot accuracy. UF IFEval is strict-prompt accuracy; SafeRLHF means and sample standard deviations use three seeds, with its IFEval variant still requiring verification.

| Method                                     | ARC-C  | HellaSwag | MMLU   | IFEval            | GSM8K  | HB↓    |
|--------------------------------------------|--------|-----------|--------|-------------------|--------|--------|
| <i>UltraFeedback: one policy seed</i>      |        |           |        |                   |        |        |
| Base                                       | 0.6049 | 0.8008    | 0.6823 | 0.7542            | 0.8666 | 0.2812 |
| NBPO                                       | 0.6058 | 0.7984    | 0.6816 | 0.7227            | 0.8605 | —      |
| Game-utilitarian                           | 0.6032 | 0.7978    | 0.6817 | 0.7209            | 0.8658 | —      |
| <i>SafeRLHF: reported three-seed study</i> |        |           |        |                   |        |        |
| Base                                       | —      | —         | —      | 0.752             | 0.867  | 0.281  |
| NBPO                                       | —      | —         | —      | $0.748 \pm 0.013$ | —      | —      |
| Game-utilitarian                           | —      | —         | —      | $0.745 \pm 0.025$ | —      | —      |

**Teacher-defined and independent evaluation.** UltraFeedback policy evaluation uses the frozen 2000-prompt final split and the same cached base response for every arm. A separate local judge – Qwen3-14B at revision 40c06982, a different model family from the ModernBERT teacher – evaluates four rubrics in both presentation orders at temperature 0, with ties worth 0.5, for 16,000 scheduled judgments per arm. The supplied main-table snapshot contains 25 evaluated policy arms, with method-specific seed counts in Table 1. Under the completeness rule this panel uses – all four attributes parsed in both presentation orders – each arm retains 1,884–1,941 prompts and their intersection is 1,236. The earlier counts of 1967 and 1960 came from a different completeness rule and are superseded rather than reported alongside. Numerical whole-prompt paired intervals are now supplied for every cell, with the minimum objective recomputed inside each replicate and policy-seed spread kept separate from prompt uncertainty.

SafeRLHF surplus uses the frozen training preference models on held-out prompts, common comparator pools, and the same definition of  $\hat{d}$ . Greedy responses define point-mass decoders; they are not unbiased samples of the trained stochastic policy. Their positive surplus does not certify Algorithm 1 acceptance. Any stochastic-policy evaluation must sample both the trained policy and reference under a common declared protocol and account for prompt uncertainty.

IFEval, GSM8K, harmful-completion checks, and benign-refusal evaluation measure complementary behavior. AlpacaEval and Arena-Hard prompt sets are scored by a local reward model; these are not official leaderboard scores. Table 31 preserves the execution report’s Arena aggregate and token medians, whose subset and prompt scope remain unspecified. They do not populate subset-specific main-table cells. The exact IFEval strict metric and numerical confidence intervals must be reconciled with evaluator outputs. A confidence interval containing 0.5 does not establish win-rate equivalence.

**Refusal concentration and teacher truncation.** The audit distinguishes identical-string responses (247/1000), short responses (450/1000), refusal-prefixed responses (621/1000), and refusals among short responses (448/450). These output-defined categories overlap and are not interchangeable. Gains in a refusal subgroup do not identify a Nash-specific mechanism.

Teacher truncation affects  $26,790/84,000 = 31.89\%$  of candidate occurrences; prompt truncation affects one occurrence. Among truncated responses, the reported median loss is 232 tokens and the observed response fraction is 61%. The latter quantities must retain their aggregation definitions. This is a response-tail visibility issue for the implemented teacher. The UltraFeedback panel therefore uses a natively supported longer-context encoder and a separate full-response evaluator, rather than

Table 30: Historical single-checkpoint objective evaluations. Short-run objective entries retain their reported precision. These rows do not replace the completed three-seed comparison in Table 22.

| Method                      | Updates | $s_{\text{help}} \uparrow$ | $s_{\text{safe}} \uparrow$ | $\min_k s_k \uparrow$ |
|-----------------------------|---------|----------------------------|----------------------------|-----------------------|
| Finite-pool target $p^*$    | —       | +0.1653                    | +0.1662                    | +0.1653               |
| Base (greedy)               | —       | −0.0405                    | −0.0320                    | −0.0405               |
| NBPO-MSE (1750)             | 1750    | +0.1337                    | +0.1337                    | +0.1337               |
| NBPO-WBC (1750)             | 1750    | +0.0738                    | +0.0808                    | +0.0738               |
| NBPO-MSE (250) <sup>†</sup> | 250     | +0.119                     | +0.125                     | +0.1186               |
| NBPO-WBC (250) <sup>†</sup> | 250     | —                          | —                          | +0.0569               |
| Game-utilitarian            | 250     | —                          | —                          | —                     |
| Fixed-reference Nash        | 250     | —                          | —                          | —                     |
| BT-RM-Nash                  | 250     | —                          | —                          | —                     |
| PROSPER                     | —       | —                          | —                          | —                     |
| MOPO                        | —       | —                          | —                          | —                     |
| Game-maxmin                 | 250     | —                          | —                          | —                     |
| Scalarized DPO              | —       | —                          | —                          | —                     |

Table 31: Completed downstream snapshot supplied by the execution report. Arena is an unassigned local-proxy aggregate; token medians have an unassigned prompt scope. These values are preserved without relabeling them as hard/creative or Alpaca-specific results. This historical snapshot contains one checkpoint per row; it is not the new three-seed aggregate.

| Policy                 | $\min_k s_k$ | IFEval | GSM8K | Harm↓ | Alpaca | Arena (unassigned) | Tokens (unassigned) |
|------------------------|--------------|--------|-------|-------|--------|--------------------|---------------------|
| Base                   | −0.0405      | .752   | .867  | .281  | .500   | .500               | 457                 |
| MSE (1750)             | +0.1337      | .758   | .864  | .284  | .419   | .435               | 481                 |
| WBC (1750)             | +0.0738      | .741   | .852  | .256  | .578   | .594               | 457                 |
| MSE (250) <sup>†</sup> | +0.1186      | .760   | .870  | .269  | .497   | .516               | 460                 |
| WBC (250) <sup>†</sup> | +0.0569      | .741   | .867  | .266  | .486   | .542               | 463                 |

increasing RoBERTa’s length limit. Reported joint-positive rates 0.725/0.730 are archived with their unit unresolved; they are not population IR certificates.

## G.2 INDEPENDENT CROSS-PLAY AND FINITE-BANK DIAGNOSTICS

Within each panel, fix the evaluator, rubric wording, parsing, response sampling, and held-out prompt IDs before comparisons. The comparator bank contains the base, NBPO, fixed-reference Nash, and one strong competitor selected on development data. Use fresh responses outside the training pool and both presentation orders. Comparator selection does not access final outcomes. Table 33 records each rubric separately; six unordered comparisons per rubric supply the off-diagonal entries, with reversed cells obtained by antisymmetry and the diagonal fixed at 0.5 by convention rather than measured self-play.

For a generated policy pair, let  $\hat{J}_k(\pi, q) = \hat{W}_k(\pi, q) - 1/2$ . A complementary policy-bank diagnostic is

$$\hat{V}_k^{\mathcal{B}}(\pi) = \min_{w \in \Delta(\mathcal{B})} \left\{ \sum_{q \in \mathcal{B}} w_q \hat{J}_k(\pi, q) + \beta_k \text{KL}(w \| a) \right\}, \quad \hat{s}_k^{\mathcal{B}}(\pi) = \hat{V}_k^{\mathcal{B}}(\pi) - \hat{V}_k^{\mathcal{B}}(\pi_{\text{ref}}).$$

Here  $a$  is a full-support distribution over the comparator policies, fixed on development data and shared by every evaluated policy;  $\beta_k$  is also fixed within the panel. Record both before evaluating this diagnostic. This is a regularized finite-policy-bank game. Its KL is over comparator weights and is not the response-distribution KL in (6); consequently it is not a direct estimate of the population value  $V_k$  or population exploitability. Training-teacher values are reported separately.

The evaluator is the campaign’s declared judge configuration: the local Qwen3-14B at revision 40c06982 with thinking disabled, temperature 0, a 256-token verdict budget, both presentation orders, and ties worth 0.5. It parses 99.2–99.8% of the 4,000 judgments in each pair, matching the 99.38–99.58% of the final-split panel. An earlier set of runs left thinking enabled, which spent

Table 32: Evaluation contracts and fields to complete from actual artifacts. These counts specify the requested full evaluation sets, not a claim that every job has completed.

| Set / target count             | Metric and completion contract                                                                                                       |
|--------------------------------|--------------------------------------------------------------------------------------------------------------------------------------|
| SafeRLHF / 1000                | Teacher-defined objective values; comparator and disagreement hashes; greedy and stochastic evaluations kept separate.               |
| IFEval / 541                   | Official strict prompt and instruction accuracy; loose scores secondary. Verify the reported strict column against evaluator output. |
| GSM8K / 1319                   | Fixed zero-shot-CoT prompt and answer extractor; exact match and parse-failure count.                                                |
| HarmBench / text test          | Report actual category counts and classifier version. Use category-specific hashing where required.                                  |
| XSTest / 250 safe, 200 unsafe  | Benign refusal and unsafe-request handling separately; report classifier or blinded manual assessment.                               |
| Alpaca / 805                   | Same-prompt local reward-model win rate versus base, ties half; response length from the same prompt set.                            |
| Arena / 500 hard, 250 creative | Report subset-specific local proxy win rates; an unassigned aggregate does not populate these columns.                               |
| MT-Bench / 80                  | Optional two-turn responses; no official score without the stated judging protocol.                                                  |

Table 33: Independent cross-play: row-policy win rates against column comparators, separately for each UF rubric, on the held-out cross-play 500. The competitor is selected on development data before any cross-play response exists. Only the unordered pair is judged and the reverse cell is  $1 - W_k$ ; the diagonal follows from that convention and is not self-play. Dashes are unmeasured.

| Rubric  | Policy                                 | Base   | NBPO   | Fixed-ref. | Competitor |
|---------|----------------------------------------|--------|--------|------------|------------|
| IF      | Base                                   | 0.5000 | 0.4770 | 0.4875     | 0.4810     |
|         | NBPO                                   | 0.5230 | 0.5000 | 0.5140     | 0.4980     |
|         | Fixed-reference Nash                   | 0.5125 | 0.4860 | 0.5000     | 0.4975     |
|         | Dev-selected competitor (util_mse_s42) | 0.5190 | 0.5020 | 0.5025     | 0.5000     |
| Truth   | Base                                   | 0.5000 | 0.4860 | 0.5105     | 0.5020     |
|         | NBPO                                   | 0.5140 | 0.5000 | 0.5135     | 0.5035     |
|         | Fixed-reference Nash                   | 0.4895 | 0.4865 | 0.5000     | 0.4860     |
|         | Dev-selected competitor (util_mse_s42) | 0.4980 | 0.4965 | 0.5140     | 0.5000     |
| Honesty | Base                                   | 0.5000 | 0.4970 | 0.5055     | 0.4885     |
|         | NBPO                                   | 0.5030 | 0.5000 | 0.5175     | 0.5140     |
|         | Fixed-reference Nash                   | 0.4945 | 0.4825 | 0.5000     | 0.4935     |
|         | Dev-selected competitor (util_mse_s42) | 0.5115 | 0.4860 | 0.5065     | 0.5000     |
| Help    | Base                                   | 0.5000 | 0.4405 | 0.4530     | 0.4500     |
|         | NBPO                                   | 0.5595 | 0.5000 | 0.5085     | 0.5005     |
|         | Fixed-reference Nash                   | 0.5470 | 0.4915 | 0.5000     | 0.4905     |
|         | Dev-selected competitor (util_mse_s42) | 0.5500 | 0.4995 | 0.5095     | 0.5000     |

the verdict budget on a reasoning block: those runs parsed 2.6–6.9% at the same budget and 88.0–91.1% at 2,048 tokens, and they are retained as set-aside artifacts rather than reported, because their missingness is not ignorable: re-judging the same twelve cells under the declared configuration moves the win rates by  $-0.019$  to  $+0.025$  with a rubric-dependent sign. Use a declared common valid subset for the cross-play summary and report coverage by method, rubric, and presentation order. Invalid judgments remain missing. Whole-prompt paired resampling keeps the policy pairs and rubrics together, recomputes each minimum and the finite-bank solve, and reports numerical intervals. Training-seed variation is reported separately. The historical 1967/1960 UF counts are not an exact paired denominator. Table 1 uses its stated common intersection; the supplied cross-play summary uses the separate denominators described in its caption and awaits a synchronized recomputation. Any new analysis choices retain the history of already observed final results.

### G.3 FEEDBACK BUDGETS AND PROVENANCE

We record unique semantic comparisons separately from order-swapped calls, generated tokens, preference-model fitting/scoring, solver time, neural training, and tuning trials. Numerical iterations reuse cached matrices and do not supply new labels. Budget-matched claims require a shared feedback contract and total-compute accounting. All responses, derived pairs, and reversed orders from one source prompt remain in the same split. Manifests identify feedback source, teacher/rubric revisions, objective order, split, comparator construction, and decoding. Historical pair-split, precursor, or differently judged results are not pooled with the repaired panel. Instruction-following error types,

length, and refusal patterns are descriptive diagnostics; output-defined subgroups alone do not identify a causal mechanism.

The source reports 11.17–11.20 GPU-hours of trainer time per completed mechanism seed, 12.25 including loading/checkpointing, and 10.1 additional GPU-hours for two failed NBPO seed-42 attempts. Shared teacher preparation is reported as 1.8 GPU-hours plus 7.3 CPU-hours; target solving adds 1.2 CPU-hours. Generation/evaluation spans 0.90–1.56 GPU-hours per measured arm. Panel-level costs, excluded from per-policy totals, are 12.6 GPU-hours for capability, 1.2 for the three general-instruction evaluations, 3.3 for cross-play and 0.9 for competitor selection. These source-reported entries require stage-ledger attribution; they are not a claim of equal total compute across algorithms.

The nominal  $1.12 \times 10^6 = 4 \times 10,000 \times \binom{8}{2}$  entries count derived learner-pair/rubric slots. They must not be equated with new human labels or unique teacher comparisons. The learner–reference block has  $4 \times 10,000 \times 8 \times 8$  rubric entries, with reference–reference scoring and preference-model training accounted for separately: the cross block is 2,560,000 teacher entries on the training split and 256,000 on the development split, the reference block another 2,560,000, and the GPT-4-generated UltraFeedback supervision behind the teacher is 318,729 pointwise UltraFeedback ratings over 80,000 completions on the disjoint 20,000-prompt fitting split plus 31,889 on its 2,000-prompt development split, honesty being the only objective with missing ratings (78,729 of 80,000). A semantic comparison is identified by the prompt, actual response IDs and objective; swapped-order requests are additional calls. Reconcile these counts from manifests before claiming a matched feedback budget. The source reports shared pool hashes across targets. DPO retains 87.9–90.7% of candidate pairs under its standardized 0.05 tie threshold; MOPO consumes candidate rather than pair rows. These are different optimization exposures.

Because learner and reference occurrences are independent, an  $8 \times 8$  cross block need not have a zero diagonal or satisfy  $A_{ij} = -A_{ji}$ . The previously reported index-wise residuals do not establish general-preference structure. Verify  $A(y, z) = -A(z, y)$  using the same response identities and check actual self-pairs separately. No automatic symmetrization of cross-block indices is justified. Both checks are now measured on the training tensor. Temperature-1 sampling repeats a learner text in the comparator pool on 144 of the 10,000 prompts, giving 2,166 *actual* self-pairs; their margins are exactly 0, so  $P(y, y) = \frac{1}{2}$  holds by measurement and not only by convention. Those repeats also make the reverse direction available: the same text scored against the other comparators appears once in the cross block and once in the reference block, and across 15,162 such identity-matched comparisons the two agree to  $8.9 \times 10^{-8}$ , within float32 rounding. Since the reference block is exactly antisymmetric with an exactly zero diagonal, this establishes  $A(y, z) = -A(z, y)$  on those identities and shows the teacher score depends on the response texts rather than on which block evaluated them. By contrast the cross block’s index-wise residuals are large,  $|A + A^\top|$  averaging 0.1925 and  $|\text{diag}|$  averaging 0.1289, which is what independent pools predict and is not a defect. Final evaluation uses 16,000 rubric-order calls per arm on 8,000 semantic rubric comparisons, separate from training feedback.

Table 34: Source-reported UF-4 costs in GPU-hours. Teacher cost is shared; Total includes the per-policy training and generation/evaluation entries but excludes shared and panel-level costs. Pair/rubric slots are derived fit-data counts, not independently elicited labels. Dashes are unmeasured.

| Method                | Pair/rubric slots  | Teacher | Policy | Gen./eval. | Total | Trials |
|-----------------------|--------------------|---------|--------|------------|-------|--------|
| NBPO                  | $1.12 \times 10^6$ | 1.8     | 11.2   | 1.0        | 12.2  | 0      |
| Fixed-reference Nash  | $1.12 \times 10^6$ | 1.8     | 11.2   | 0.9        | 12.1  | 0      |
| Game-utilitarian      | $1.12 \times 10^6$ | 1.8     | 11.2   | 1.1        | 12.3  | 0      |
| BT-RM–Nash            | $1.12 \times 10^6$ | 1.8     | 11.2   | 0.9        | 12.1  | 0      |
| Global game-maxmin    | $1.12 \times 10^6$ | 1.8     | 11.2   | 0.9        | 12.1  | 0      |
| PROSPER (adapt.)      | $1.12 \times 10^6$ | 1.8     | 11.2   | —          | —     | —      |
| MOPO (adapt.)         | $1.12 \times 10^6$ | 1.8     | 6.1    | —          | —     | —      |
| DPO (uniform / sweep) | $1.12 \times 10^6$ | 1.8     | 8.4    | —          | —     | 0      |

Report generated/scored token counts, order-swapped and repeated judgments, solver CPU time, and shared-cache attribution alongside Table 34. The DPO deployment-policy cost and total seven-weight search cost are distinct. Existing component measurements do not establish an equal-total-budget comparison or a performance-versus-compute curve.

---

## G.4 EMPIRICAL PARETO EVALUATION

The population Nash optimum is Pareto efficient under the stated assumptions; a single proximal step and approximate neural realization do not inherit this certificate. The instruction comparison evaluates all four attributes using a fixed independent local judge; the measured methods and seed counts are given in Table 1. A policy is empirically nondominated only relative to the evaluated policies and evaluator in this four-dimensional space. The two projections in Figure 2 cannot establish four-dimensional dominance.

Scalarized DPO (Rafailov et al., 2023) uses development-calibrated scalar BT rewards. Its declared weights are the four coordinate vectors,  $(.25, .25, .25, .25)$ ,  $(.1, .1, .1, .7)$  and  $(.1, .7, .1, .1)$  in IF/Truth/Honesty/Help order. The main table represents the uniform setting. The supplied source reports 3/7 weight settings in progress or completed. Only the uniform setting is reconstructed in the main figure; the complete sweep requires its original artifacts.

The reported MOPO adaptation fixes IF as primary and constrains the other objectives. It uses the stated adaptive-bound parameters  $\beta = 0.9995$ ,  $\tau = 0.08$ , and the main setting  $t_0 = \infty$ ; a  $t_0 = 250$  variant must be labeled separately. The original robustification radius  $\varepsilon = 0.15$  produced reference constraint deficits of 0.061–0.083, versus nominal slack  $2.5 \times 10^{-4}$ , and an increasing multiplier (73.8 after 300 iterations). The reported run therefore removes that robustification and applies plug-in constraints. Exact evaluation of a cached tensor removes numerical integration error on this pool, not preference-model, sampling or generalization error. This is a MOPO-inspired plug-in adaptation, not a faithful reproduction of the robust algorithm (Agnihotri et al., 2026). It also uses the shared full-parameter training budget rather than the original LoRA recipe.

In the reported six-stage target simulation, secondary multipliers remain zero while bounds rise from 0.500 to 0.570/0.592/0.603. IF–secondary candidate-rank correlations are 0.870/0.912/0.905, positive on 99.6–100% of prompts. The IF-best candidate exceeds each secondary’s pool mean on 99.0–100% of prompts. Thus the constraints are inactive in this cached-pool simulation, suggesting limited conflict under this teacher. This does not establish agreement under independent feedback or explain neural outcomes causally. Table 1 retains the measured DPO, MOPO and PROSPER rows with their actual seed counts.

Compare methods on the complete final set and prespecified label-disagreement subsets. Keep independent win rates separate from training-teacher game values. Report individual policy seeds, paired prompt uncertainty, and all declared weights. Curves joining points do not represent evaluated mixture policies. Nash welfare is defined only for positive game surpluses; mark other evaluated policies outside the positive-surplus domain. Fix hypervolume normalization and reference points before final scores are opened.

**Uncertainty and provenance.** Report every policy seed, then mean and standard deviation across training seeds. Each paired prompt-bootstrap replicate recomputes objective means and then the minimum. Its intervals condition on a checkpoint and do not include training-seed variability. Freeze development choices before confirmatory comparisons and retain the history of benchmark exposure. SafeRLHF Nash/utilitarian aggregates and UF neural diagnostics use three policy seeds; UF generated-policy seed counts vary by method as reported in Table 1; historical WBC/long-schedule snapshots use one. Missing method-wise values must not be reconstructed from ranges pooled across methods; dashes are unmeasured, not zeros. All result updates must record checkpoint, split, decoding, evaluator, sample count, and source artifact. Unmatched literature scores do not fill matched-baseline rows.

## H TARGET-TO-POLICY DIAGNOSTIC PROTOCOL

This development-only diagnostic leaves completed UF-4 results unchanged. Fix 200 policy-dev prompts and 200 disjoint training prompts by an outcome-independent rule. Record IDs and model, tokenizer, template and decoding revisions. A reduced budget is a separately declared pilot, not a silent replacement.

**Independent scoring on common support.** Retain the eight learner occurrences  $Y_x$ , frozen targets and shared training multipliers; do not refit a 200-prompt bargain. Select four reference occurrences  $Z_x^E$  independently of outcomes, preserving duplicate multiplicity. Independent evaluation supplies

$$A_{k,x}^E[i, j] = \widehat{P}_k^E(y_{x,i} \succ z_{x,j}^E | x) - \frac{1}{2}, \quad B_{k,x}^E[\ell, j] = \widehat{P}_k^E(z_{x,\ell}^E \succ z_{x,j}^E | x) - \frac{1}{2}.$$

Average both orders for each reference pair, set actual self-pairs to zero and reconstruct reversed pairs by antisymmetry. Cross-matrix index transposition is not response reversal. Evaluation scores never refit targets or select weights.

With  $\mu_j^E = 1/4$  and one common evaluation temperature  $\beta_k^E = 0.25$ , define

$$W_k^E(p) = \mathbb{E}_x \left[ \frac{1}{2} + \sum_{i,j} p_{x,i} \mu_j^E A_{k,x}^E[i, j] \right], \quad (44)$$

$$V_{k,x}^E(p) = -\beta_k^E \log \sum_j \mu_j^E \exp \left\{ -\frac{\sum_i p_{x,i} A_{k,x}^E[i, j]}{\beta_k^E} \right\}, \quad (45)$$

$$d_{k,x}^E = -\beta_k^E \log \sum_j \mu_j^E \exp \left\{ -\frac{\sum_\ell \mu_\ell^E B_{k,x}^E[\ell, j]}{\beta_k^E} \right\}, \quad s_k^E(p) = \mathbb{E}_x [V_{k,x}^E(p) - d_{k,x}^E]. \quad (46)$$

These are finite-matrix evaluations, not unbiased population game-value estimates. Reference wins and surplus are .5 and 0 by convention. The empirical learner distribution  $p_{t,x}$  is different and must not receive these values by definition. Training temperatures and targets remain unchanged.

**Fitted pool versus generated policy.** For each frozen full-run checkpoint, evaluate

$$\tilde{\theta}_{\theta,x}(i) = \frac{p_{t,x}(i) \exp\{\log \pi_\theta(y_{x,i} | x) - \log \pi_t(y_{x,i} | x)\}}{\sum_{i'} p_{t,x}(i') \exp\{\log \pi_\theta(y_{x,i'} | x) - \log \pi_t(y_{x,i'} | x)\}}.$$

Sum log probabilities over identical response tokens and terminal-token conventions. Raw-likelihood normalization is not this diagnostic. Fresh rows compare one new stochastic response per prompt with the same four references under common decoding. Report direct wins; one draw does not estimate the nonlinear stochastic-policy game value, so its surplus is n/a.

**Coverage and interpretation.** Report planned/common-complete counts and invalid rates by rubric and order. High per-call parsing accuracy need not imply complete prompts. After uniform repair, retain unresolved entries as missing. Bound their probabilities in  $[0, 1]$ ; combine lower policy and upper reference values for a surplus lower bound, and conversely for its upper bound. Complete-case scores are conditional estimates. Use paired difference intervals; containing zero does not establish equivalence.

**Equivalent all-candidate regression.** On a prompt with  $n$  learner occurrences define

$$u_i = \log \frac{\pi_\theta(y_{x,i} | x)}{\pi_t(y_{x,i} | x)} - \log \frac{p_{t,x}^*(i)}{p_{t,x}(i)}, \quad \bar{u} = \frac{1}{n} \sum_i u_i.$$

For uniform unordered pairs, Eq. (26) can be evaluated exactly as

$$\frac{1}{\binom{n}{2}} \sum_{i < i'} (u_i - u_{i'})^2 = \frac{2}{n-1} \sum_i (u_i - \bar{u})^2. \quad (47)$$

The identity follows from  $\sum_{i < i'} (u_i - u_{i'})^2 = n \sum_i (u_i - \bar{u})^2$  and needs one likelihood per candidate. Nonuniform pair measures require their original weights. Verify values and gradients against explicit pairs. Both full-parameter pilot arms share initialization, targets and predeclared learning rate, token budget and stopping rule. Report train/dev fit,  $\text{KL}(p^* \|\tilde{p}_\theta)$ , lengths and fresh wins; keep these pilots separate from the 10k-prompt, 1,250-update runs.

Table 35: Fixed diagnostic contract, with the scope actually completed. A call containing several rubrics must be distinguished from the number of rubric-level judgments it returns.

| Component              | Contract                                                                                                                                                                                                                                                                                                        |
|------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Prompts and pools      | 200 policy-dev prompts; 8 learner and 4 reference occurrences; all categorical rows reuse the same matrices.                                                                                                                                                                                                    |
| Evaluator              | Campaign’s verified independent Qwen3-14B revision and rubric; thinking off; both orders; ties .5. Cache and log the actual configuration.                                                                                                                                                                      |
| Categorical workload   | $200(8 \times 4 + \binom{4}{2}) \times 4 \times 2 = 60,800$ rubric-order judgments before repairs.                                                                                                                                                                                                              |
| Fresh workload         | $200 \times 4 \times 4 \times 2 = 6,400$ rubric-order judgments per policy, plus 200 generated responses.                                                                                                                                                                                                       |
| Uncertainty            | 2,000 whole-prompt paired bootstrap replicates; recompute objective means, minima, finite-bank surpluses and method differences. Seed SD remains separate.                                                                                                                                                      |
| Missing entries        | Uniform predeclared formatting repair for every invalid verdict; retain original and repaired outputs. No tie imputation or renormalization around missing cells.                                                                                                                                               |
| Completed, categorical | 60,800 of 60,800 judged on all 200 prompts; 60,252 parsed (99.10%); complete prompts per objective 143/131/144/167 and 78 common to all four; $[0, 1]$ bounds computed on all 200.                                                                                                                              |
| Completed, fresh       | Six arms at 6,400 each, 38,400 judgments on 1,200 generated responses: NBPO, utilitarian and uniform DPO at seed 42, both projection pilots, and the optional untrained-base calibration draw.                                                                                                                  |
| Completed, pilots      | Two full-parameter arms, 200 optimizer steps and 6,597,080 non-padding response tokens each on 200 training prompts, 2.00 GPU-hours together.                                                                                                                                                                   |
| Evaluator as run       | Local Qwen3-14B, revision recorded as 40c06982, rubric file hash 7f6af77e, thinking off, temperature 0, 256 new tokens, both orders, ties .5; no paid API call. The directory carries no upstream snapshot identifier, so the revision is our recorded declaration rather than an independently verifiable tag. |
| Evaluation constants   | $\beta^E = 0.25$ and $m_{\text{eval}} = 4$ fixed before the panel ran; bootstrap seed 20260914 with 2,000 replicates; panel hash f899600a; training panel hash 02a1c2ce.                                                                                                                                        |

**Target separation and objective conflict.** Recompute all entries in Table 36 on this same development panel and pair measure. Target-difference RMS is  $\sqrt{\mathbb{E}[(\hat{h}_a^* - \hat{h}_b^*)^2]}$ . Let  $e_a = h_a - \hat{h}_a^*$ ; fit-error RMS is  $\sqrt{\mathbb{E}[e_a^2]}$  and error-difference RMS is  $\sqrt{\mathbb{E}[(e_a - e_b)^2]}$ . The identity  $h_a - h_b = (\hat{h}_a^* - \hat{h}_b^*) + (e_a - e_b)$  distinguishes shared fitting error from error that changes the method contrast. Compare their scales without interpreting their ratio as a statistical power calculation or proof of signal loss. Report cross-objective winner disagreement, non-tie denominators and rank correlations. These measure conflict, not within-objective cycles; independent-judge agreement assesses teacher transfer.

Measured on that panel, the two terms are far apart in scale. The NBPO and utilitarian targets differ by RMS 0.179 and the NBPO and fixed-reference targets by 0.418, against a target RMS of 5.38 and mean total variations of 0.021 and 0.024 between the target distributions themselves. The fitted arms’ errors are RMS 4.21 for NBPO and 4.28–4.30 for the controls, so the regression leaves most of the target unexplained at this budget. Their error *difference* is RMS 1.18 and 1.23, which is 6.6 and 2.9 times the corresponding target difference, and the difference between the two fitted policies is RMS 1.18 and 1.20 – numerically the error difference rather than the target difference. We report these as scales and draw no power calculation from their ratio. Read with the independent panel, where the same three targets differ by at most 0.0018 in judged win rate, the two measurements say the same thing from different directions: what distinguishes these aggregation rules is small both in the numbers the regression must reproduce and in the quality an independent evaluator can detect.

The conflict report on the same panel says why the targets are so close. Over each prompt’s eight candidates the four objectives are strongly co-ranked: Spearman correlations average 0.777 to 0.915 with medians of 0.857 to 0.952, at most 1.0% of prompts show a negative correlation for any pair, the two objectives agree on the single best candidate on 57.0% to 77.5% of prompts, and the candidate maximizing one objective still clears the other’s pool mean on 96.0% to 100%. Truthfulness against helpfulness is the least aligned pair on every one of those measures. Target distributions follow: mean total variation between the Nash target and the utilitarian target is 0.0206 (maximum 0.0732) and against the fixed-reference target 0.0240 (maximum 0.1281). This is between-objective conflict, and it is small. It is not a statement about cycles within an objective: on the scalar values this panel induces a strict order except for exact ties, so its weak no-Condorcet rate of 0 to 0.5% and its zero cyclic triangles over 11,200 triples are close to automatic and should not be read as evidence about the full pairwise tensor, which Appendix D.1 measures separately.

Table 36: Measured target diagnostics on the same 200 development prompts and pair measure. Target differences and fit-error differences are descriptive scales, not a causal attribution of policy performance.

| Targets compared       | <i>N</i> | Mean TV | Target diff. RMS | NBPO error RMS | Error diff. RMS |
|------------------------|----------|---------|------------------|----------------|-----------------|
| NBPO / utilitarian     | 200      | .0206   | .1794            | 4.2117         | 1.1806          |
| NBPO / fixed reference | 200      | .0240   | .4178            | 4.2117         | 1.2321          |

Also report the fraction of prompts where the NBPO target has positive surplus on every objective. This is selected-target behavior, not an infeasibility certificate. Per-prompt and aggregate positivity differ.

## I PROSPECTIVE CONFLICT AND CYCLICITY STRESS TESTS

This section is a preregistered extension template; dashes denote unmeasured quantities, not zero. It does not change the completed UF experiments. Selection uses development data and measured cost before final-test access. A negative screening result is retained.

**Separate the mechanisms.** Within-objective cycles test preference representation and opponent adaptation; cross-objective disagreement tests compromise selection. Neither guarantees a Nash advantage over other game-based methods. We first measure stable cycles, objective conflict, joint strict feasibility, and separation of exact targets. A large raw cycle count alone is insufficient: it increases with the number of candidate triples and can involve only weak responses.

**Dataset contract.** We screen fresh responses on SafeRLHF prompts under fixed helpfulness and harmless rubrics, and WildChecklists prompts under their native checklist items. PROSPER uses the latter with prompt-dependent objectives (Zhang et al., 2026). Native checklist positions are not shared objective identities and cannot be substituted for the global  $k$  in NBPO. WildChecklists fine-tuning requires a separately frozen set of semantically constant global rubrics; such feedback is an adaptation. Existing SafeRLHF pair labels support a conflict audit but do not by themselves supply complete response triangles. New LLM judgments are AI feedback, not new human annotations.

**Screening.** Use 200 development prompts per new source and eight responses per prompt. On the same 28 unordered pairs, collect both orders and two independent repetitions per order; use a preselected 50-prompt subset for five additional judgments per order. Report normalized cyclic-triangle rates, repeatability, no-Condorcet rates, order sensitivity, and a fitted-BT noise control. Estimate objective conflict on the same non-tied pairs. Report prompt-level variation rather than pooling checklist positions as global parties. Table 37 is a compact index; per-objective denominators and uncertainty remain in the accompanying result artifact.

Table 37: Development screening template.  $N$  is the common prompt count;  $C$  is the repeated cyclic-triangle fraction;  $D$  is cross-objective pair disagreement;  $\gamma^* = \max_p \min_k [V_k(p) - d_k]$  uses the specified aggregate finite problem. TV compares exact NBPO and utilitarian targets. Checklist-native feasibility and target columns are inapplicable until fixed objective identities are defined.

| Feedback contract             | $N$ | $C$    | $D$    | $\gamma^*$ | Target TV |
|-------------------------------|-----|--------|--------|------------|-----------|
| UF / fresh fixed rubrics      | —   | —      | —      | —          | —         |
| SafeRLHF / help–safe          | 199 | 0.0010 | 0.1269 | 0.2608     | 0.0069    |
| WildChecklists / native items | —   | —      | —      | n/a        | n/a       |

**Factorial mechanism control.** Cross low/high cyclicity with low/high objective conflict using matched finite games, fixed payoff scales and reference distributions. Record realized factors and the joint-feasibility margin; do not select games by method ranking. Evaluate exact targets before fitting a policy. In Table 38,  $\Delta M$  is the NBPO-minus-fixed-reference difference in normalized worst surplus under the same adaptive evaluator. TV compares NBPO with utilitarian targets. These constructed games identify mechanisms and are not evidence of naturally cyclic human preferences; no cell is required to favor NBPO.

Table 38: Controlled factorial template. All cells retain the same declared game seeds and evaluation contract;  $C$ ,  $D$  are measured factor checks. Report means and uncertainty in the result artifact.

| Cycle setting | Conflict setting | $C$    | $D$    | $\gamma^*$ | $\Delta M$ | Target TV |
|---------------|------------------|--------|--------|------------|------------|-----------|
| Low           | Low              | 0.0004 | 0.0719 | 0.0642     | 0.0000     | 0.0005    |
| Low           | High             | 0.0004 | 0.4519 | 0.0642     | 0.0006     | 0.0115    |
| High          | Low              | 0.4277 | 0.0311 | 0.0642     | 0.0123     | 0.0003    |
| High          | High             | 0.3594 | 0.1331 | 0.0642     | 0.0243     | 0.0065    |

**One confirmatory policy panel.** Select at most one compatible feedback contract, with 2,000/500/1,000 disjoint train/development/test prompts (a cost-based reduction is frozen before outcome inspection). Keep the completed UF panel as a transitive-supervision control. All new methods share source pools, feedback, base initialization, token accounting, and a declared tuning budget. The default two objectives are helpfulness and harmless; changing them requires a separately named contract. The full test distribution is primary; a conflict stratum defined before policy evaluation is secondary. Seeds and missing-data coverage cannot be silently reduced. BT-projected Nash is a projection control, not a reproduction of a globally trained reward model.

Table 39: New fixed-objective policy panel template; dashes denote unmeasured values.  $W_1, W_2$  are independent reference wins, and  $\bar{W}_{\min}$  averages seed-specific minima.  $\Delta$  is the paired difference in this statistic versus NBPO, with a 95% prompt interval in the underlying artifact. One fresh response does not identify a policy game value.

| Method                       | Seeds | $N$ | $W_1$  | $W_2$  | $\bar{W}_{\min}$ | $\Delta$ |
|------------------------------|-------|-----|--------|--------|------------------|----------|
| Base fresh draw              | 1     | 986 | 0.5011 | 0.4896 | 0.4896           | —        |
| NBPO                         | —     | —   | —      | —      | —                | —        |
| Fixed-reference Nash         | 1     | 986 | 0.6786 | 0.6334 | 0.6334           | —        |
| Game-utilitarian             | —     | —   | —      | —      | —                | —        |
| Global game-maxmin           | —     | —   | —      | —      | —                | —        |
| BT-projected Nash            | —     | —   | —      | —      | —                | —        |
| PROSPER (matched adaptation) | —     | —   | —      | —      | —                | —        |
| Scalarized DPO (uniform)     | —     | —   | —      | —      | —                | —        |

Results pending: exact target / fitted pool / fresh-response objective vectors  
Paired contrasts on common prompts; full test and prespecified conflict stratum

Figure 4: Prospective mechanism-to-policy figure. Report target and fitted-pool contrasts on identical development candidates, and fresh-policy contrasts on a separate fixed common test set. Use paired prompt intervals within each panel and display all applicable arms at each stage (undefined stages are n/a), with all predeclared arms in the fresh panel. Do not subtract development and test stages as a transfer effect. No monotone transfer or superiority is assumed.